

A Little Order!

Delving into the STL sorting algorithms

Fred Tingaud

@FredTingaudDev

quick-bench.com

Quick C++ Benchmark

More ▾

```
1 static void StringCreation(benchmark::State& state) {
2     // Code inside this loop is measured repeatedly
3     for (auto _ : state) {
4         std::string created_string("hello");
5         // Make sure the variable is not optimized away by compiler
6         benchmark::DoNotOptimize(created_string);
7     }
8 }
9 // Register the function as a benchmark
10 BENCHMARK(StringCreation);
11
12 static void StringCopy(benchmark::State& state) {
13     // Code before the loop is not measured
14     std::string x = "hello";
15     for (auto _ : state) {
16         std::string copy(x);
17     }
18 }
19 BENCHMARK(StringCopy);
20
```

compiler = clang-5.0 ▾

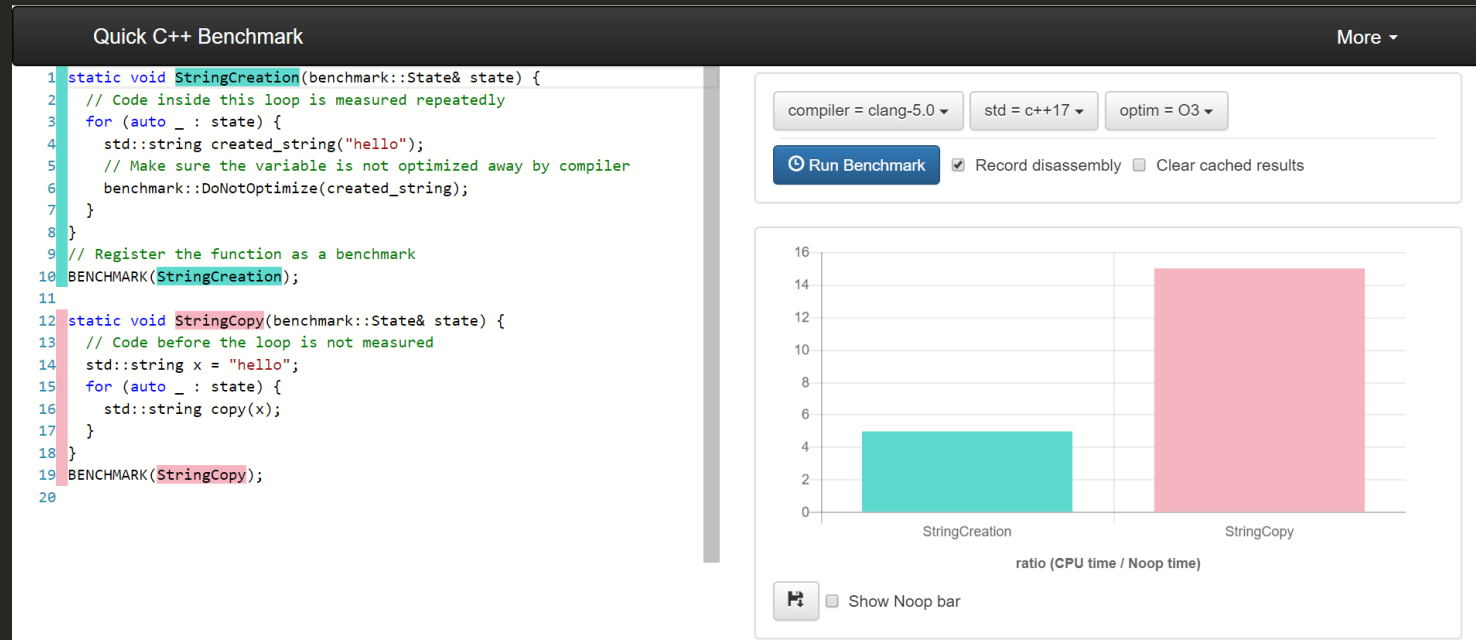
std = c++17 ▾

optim = O3 ▾

⌚ Run Benchmark

☒ Record disassembly

quick-bench.com




A use case

Sorting a 1,000,000 elements vector to get the median.

Clang 3.8


GNU's libstdcxx

std::sort




std::stable_sort

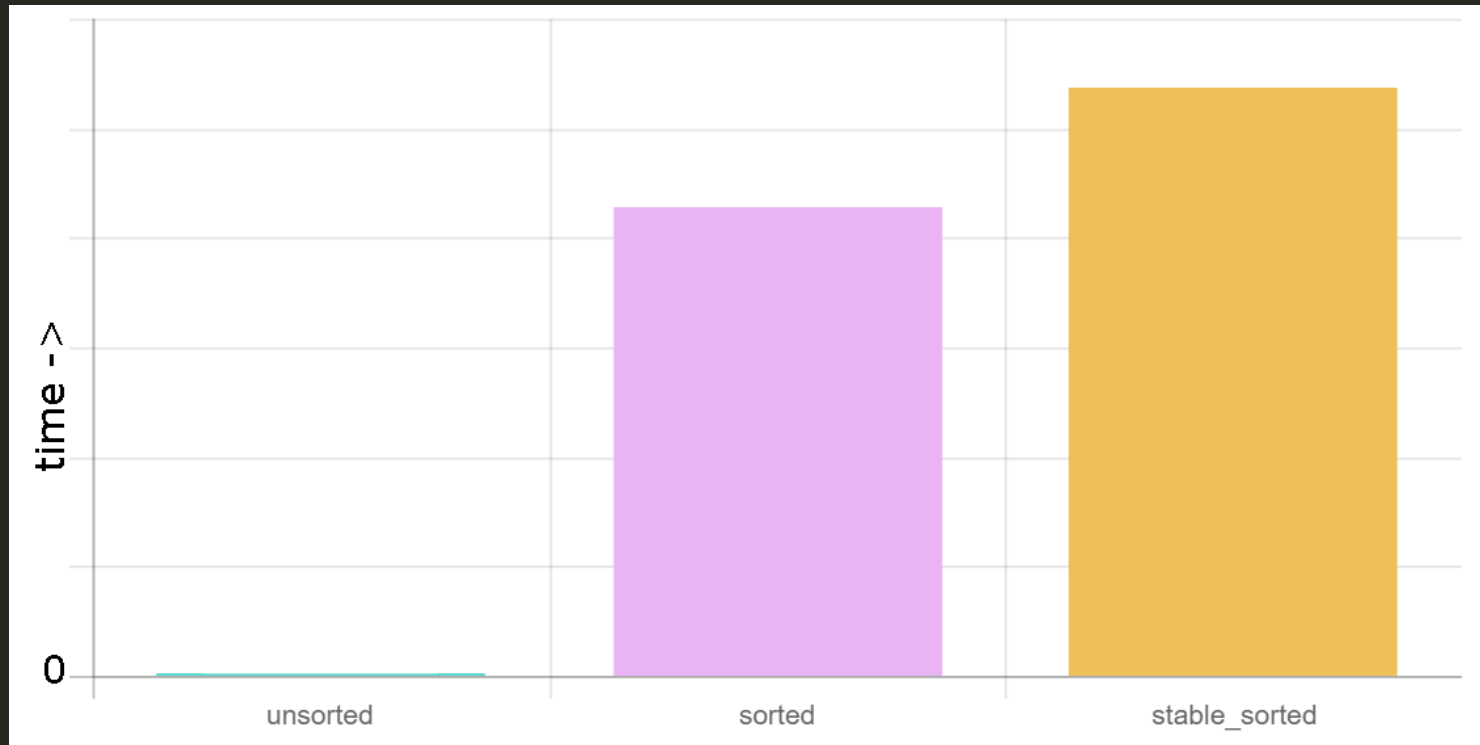
`std::sort`




`std::stable_sort`



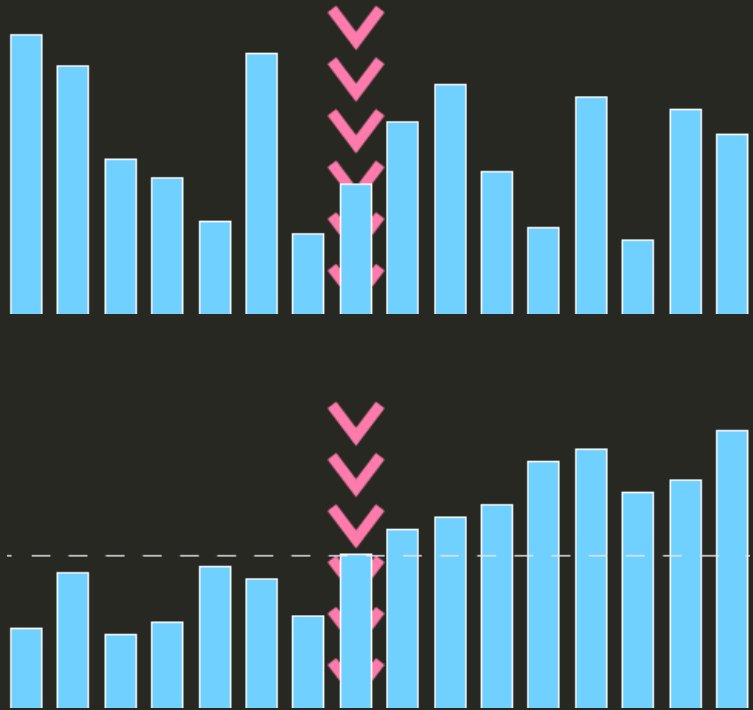
std::stable_sort



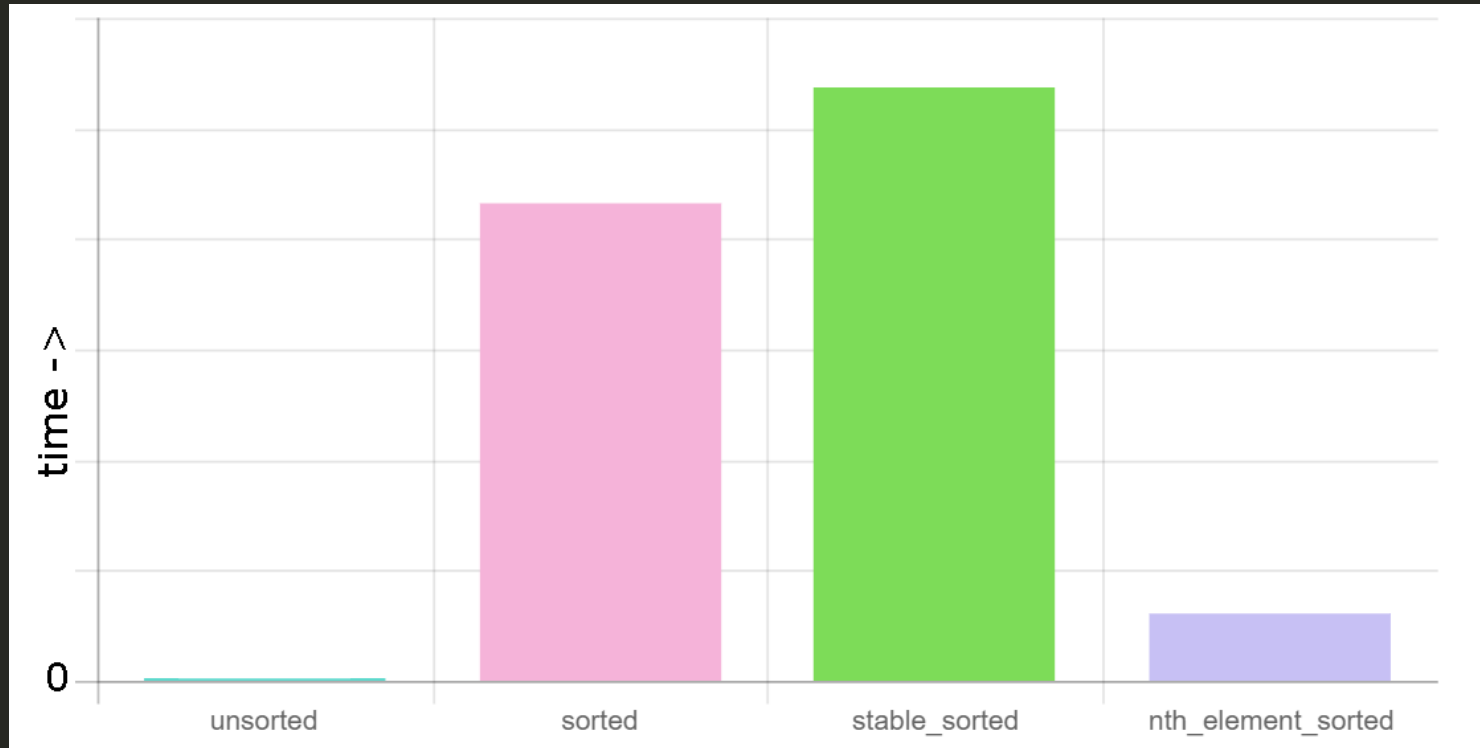
std::nth_element



std::nth_element




std::nth_element



nth_element is 10x faster than sort

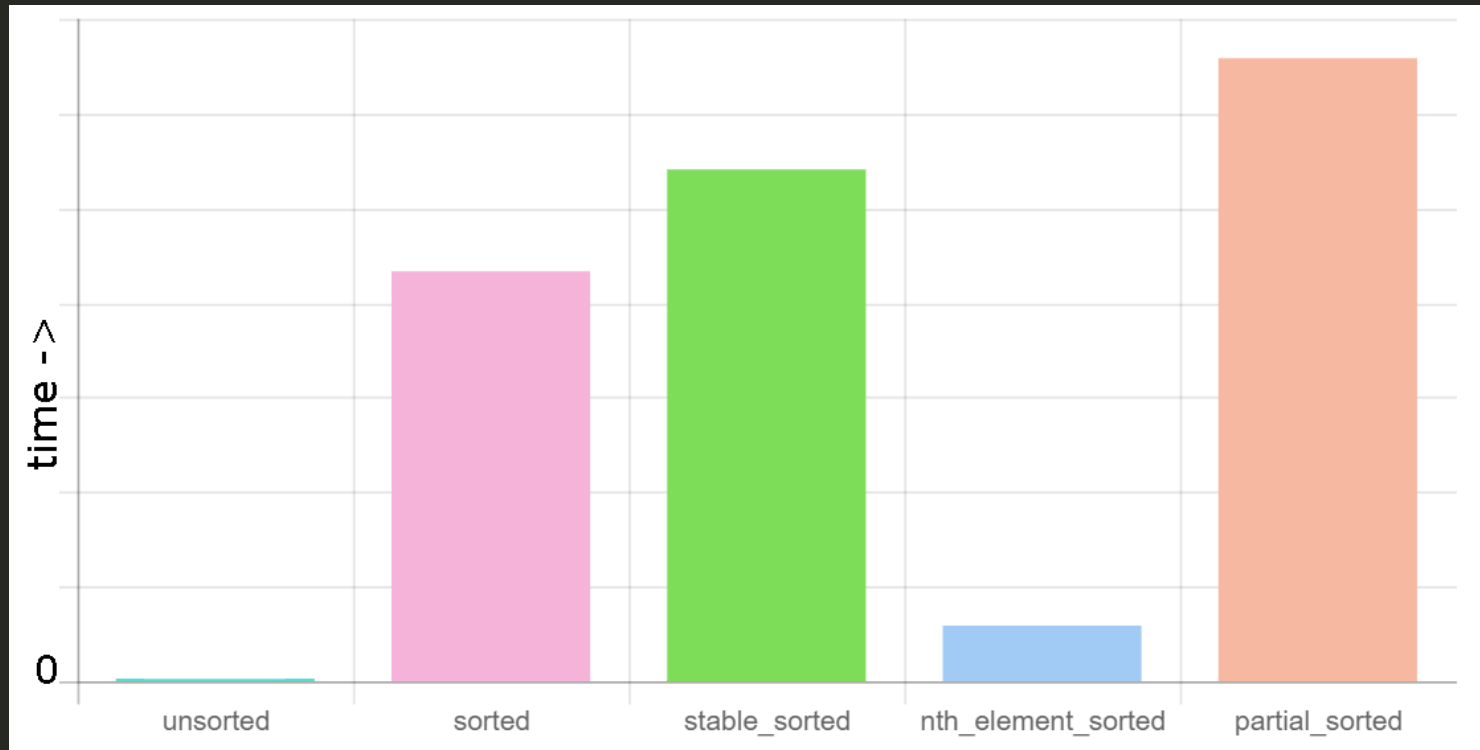
std::partial_sort



std::partial_sort



std::partial_sort




Home-cooked partial sort

What if we just call `std::sort` on the result of `std::nth_element`?

It should do the same as `std::partial_sort...`

Home-cooked partial sort



Home-cooked partial sort

Did we do better than the STL implementers in 5 minutes and 2 lines of code?

Home-cooked partial sort

Did we do better than the STL implementers in 5 minutes and 2 lines of code?

Probably not...

Home-cooked partial sort

Did we do better than the STL implementers in 5 minutes and 2 lines of code?

Probably not...

Let's try to understand why!

Check the contract!

Read the Standards

- Draft: <https://github.com/cplusplus/draft>
- C++ Reference: <https://cppreference.com>

Read the Standards - Sort

std::Sort

Defined in header <algorithm>

```
template< class RandomIt >                                (until C++20)
void sort( RandomIt first, RandomIt last );                (1)
template< class RandomIt >                                (since C++20)
constexpr void sort( RandomIt first, RandomIt last );
template< class ExecutionPolicy, class RandomIt >           (since C++17)
void sort( ExecutionPolicy&& policy, RandomIt first, RandomIt last );
template< class RandomIt, class Compare >                  (until C++20)
void sort( RandomIt first, RandomIt last, Compare comp );  (3)
template< class RandomIt, class Compare >                  (since C++20)
constexpr void sort( RandomIt first, RandomIt last, Compare comp );
template< class ExecutionPolicy, class RandomIt, class Compare > (since C++17)
void sort( ExecutionPolicy&& policy, RandomIt first, RandomIt last, Compare comp ); (4)
```

Sorts the elements in the range `[first, last)` in ascending order. The order of equal elements is not guaranteed to be preserved.

- 1) Elements are compared using operator<.
- 3) Elements are compared using the given binary comparison function comp.
- 2,4) Same as (1,3), but executed according to policy. These overloads do not participate in overload resolution unless `std::is_execution_policy_v<std::decay_t<ExecutionPolicy>>` is true

Parameters

first, last - the range of elements to sort

policy - the execution policy to use. See [execution policy](#) for details.

comp - comparison function object (i.e. an object that satisfies the requirements of *Compare*) which returns `true` if the first argument is *less* than (i.e. is ordered *before*) the second.

The signature of the comparison function should be equivalent to the following:

```
bool cmp(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const` &, but the function object must not modify the objects passed to it.

The types `Type1` and `Type2` must be such that an object of type `RandomIt` can be dereferenced and then implicitly converted to both of them.

Type requirements

- `RandomIt` must meet the requirements of *ValueSwappable* and *RandomAccessIterator*.
- The type of dereferenced `RandomIt` must meet the requirements of *MoveAssignable* and *MoveConstructible*.
- `Compare` must meet the requirements of *Compare*.

Return value

(none)

Complexity

$O(N \log N)$, where $N = \text{std::distance}(\text{first}, \text{last})$ comparisons on average. (until C++11)

$O(N \log N)$, where $N = \text{std::distance}(\text{first}, \text{last})$ comparisons. (since C++11)

Exceptions

Standards - `std::sort`

Complexity

$O(N \cdot \log(N))$, on average (until C++11).

$O(N \cdot \log(N))$ (since C++11).

Standards - `std::partial_sort`



Complexity

Approximately $(\text{last-first})\log(\text{middle-first})$
applications of `cmp`



$O(N \cdot \log(k))$

- N = full container size
- k = subset size

Two complexities



Two complexities



Read the Standards - Partial Sort

For a median, $k = N / 2$

$$O(N \cdot \log(k)) \Leftrightarrow O(N \cdot \log(N / 2))$$

Read the Standards - Partial Sort

For a median, $k = N / 2$

$O(N \cdot \log(N))$

Read the Standards - Partial Sort

For a median, $k = N / 2$

$O(N \cdot \log(N))$

Just like `std::sort`

Read the Standards - Partial Sort

If we want to sort the 10 best scores in a MMO scoreboard.

$k = 10$

$O(N \cdot \log(10))$

Read the Standards - Partial Sort

If we want to sort the 10 best scores in a MMO scoreboard.

$k = 10$

$O(N)$

Read the Standards - Nth Element + Sort

The contract of Nth Element:

$O(N)$ on average

So Nth Element + Sort on the result:

$O(N + k \cdot \log(k))$

Read the Standards - Nth Element + Sort

When $k = N / 2$

$$O(N + N/2 \cdot \log(N/2)) \Leftrightarrow O(N \cdot \log(N))$$

When $k = 10$

$$O(N + 10 \cdot \log(10)) \Leftrightarrow O(N)$$

Read the Standards


Same complexities for `std::partial_sort` and our algorithm.

Compare implementations / compilers

Comparing implementations

- Compiler Explorer: <https://godbolt.org>
- Wandbox: <https://wandbox.org>

Comparing implementations



Comparing implementations

Same results for all implementations of the libc++.

Measure

Measure


- Google Benchmark:
<https://github.com/google/benchmark>
- Quick Bench: <http://quick-bench.com>

Measure Size Variation

We are sorting a subset inside a container.

- Subset size may vary.
- Container size may vary.
- Both may vary.


Measure Size Variation - Subset



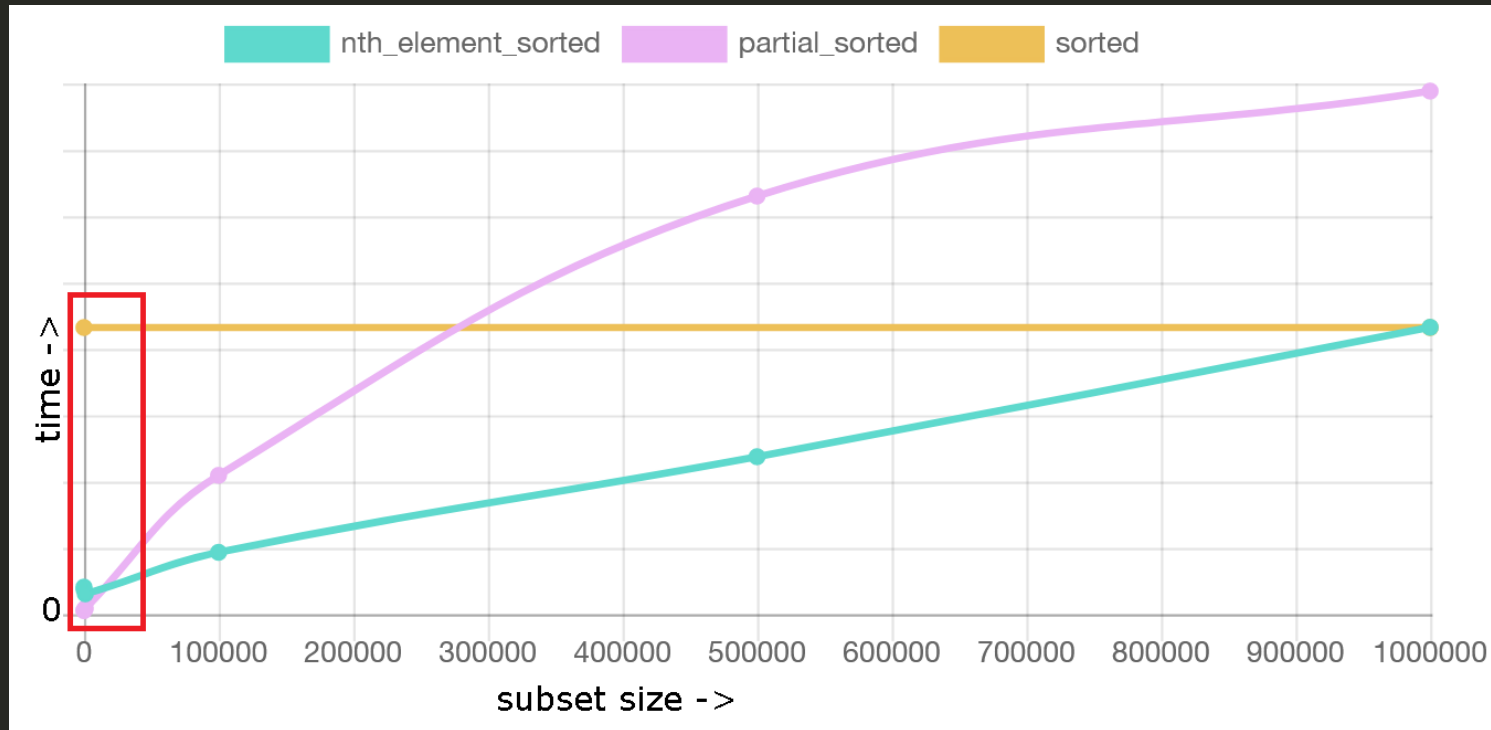
Measure Size Variation - Subset

- 1,000,000 elements
- We vary the number of elements we sort in it

Measure Size Variation - Subset

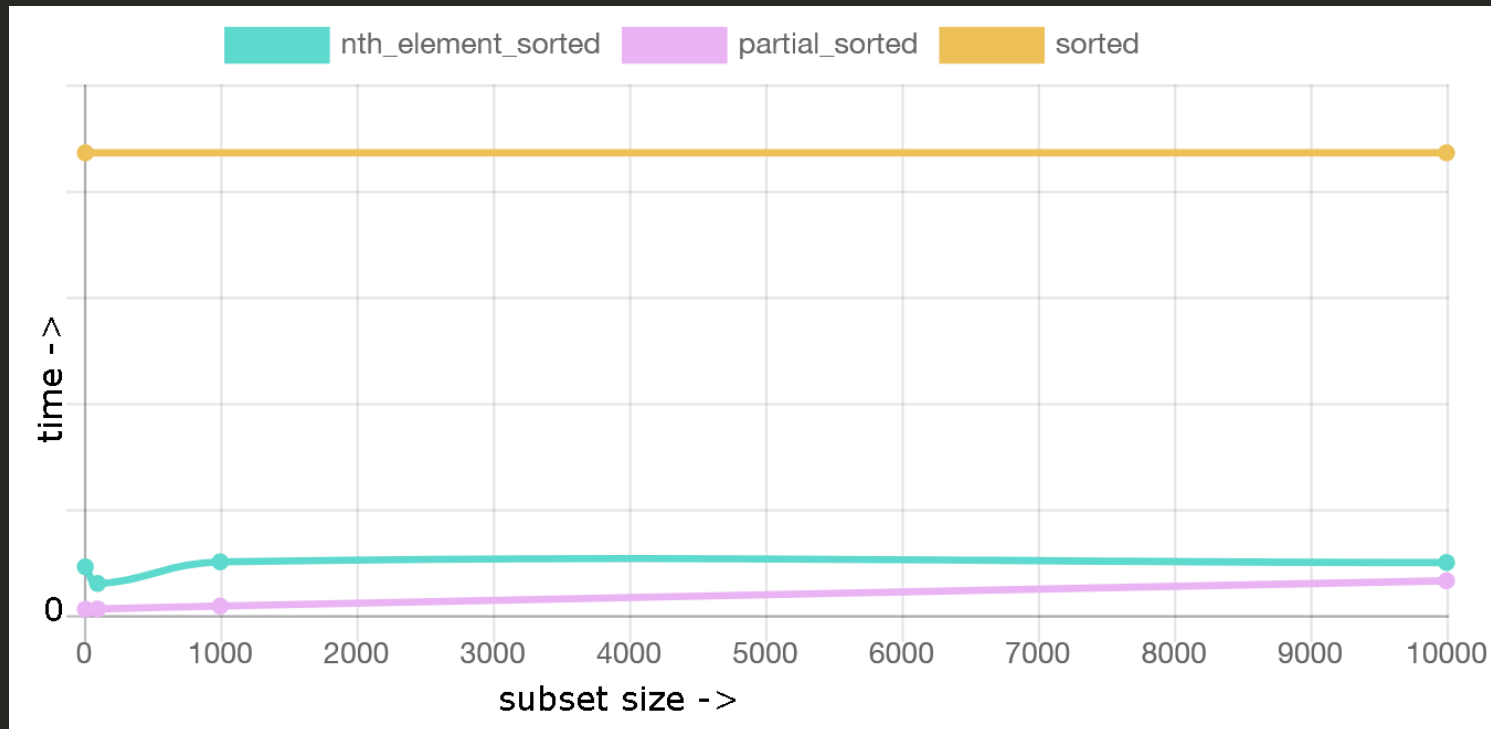


Measure Size Variation - Subset




Measure Size Variation - Subset

Zoomed



Measure Size Variation - Container

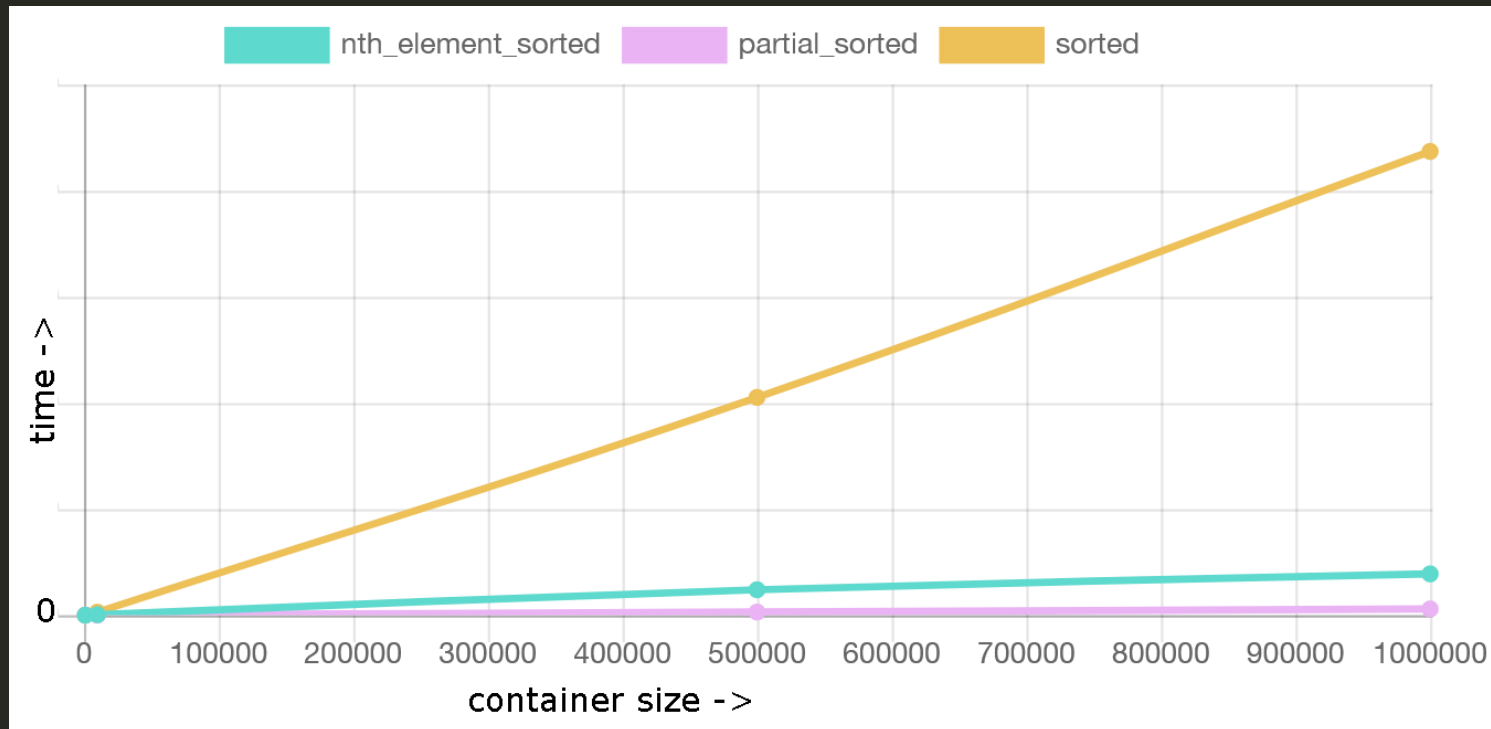


Measure Size Variation - Container


- We sort 100 elements
- We vary the size of the full container

Measure Size Variation - Container

$k = 100$



Measure Size Variation - Both




Measure Size Variation - Both

- We sort 1/5th of the elements
- We vary the size of the container

Measure Size Variation - Both

$$k = N / 5$$



Measure Size Variation

We have an actual difference!

When to use `std::partial_sort`

Use `partial_sort` when sorting a subset that is considerably smaller than the whole container.

Otherwise, use `nth_element + sort`.

Why?

How is `partial_sort` faster than `nth_element` for lower values of `k`?

Why is it slower for higher values?

Read The Code

Read The Code

- GCC - <https://github.com/gcc-mirror/gcc>
- LLVM's libcxx - <https://github.com/llvm-mirror/libcxx>
- Visual Studio - `F12`


```
std::sort
```

```
template<typename RandIt, typename _Compare>
inline void
__sort(RandIt first, RandIt last,
       _Compare comp) {
    if (first != last){
        std::__introsort_loop(first, last,
                               std::__lg(last-first)*2,
                               comp);
        std::__final_insertion_sort(first, last, comp);
    }
}
```

std::sort

Complexity

$O(N \cdot \log(N))$, on average (until C++11).

$O(N \cdot \log(N))$ (since C++11).

Introsort

- Quicksort is very fast for most scenarios, but it can become $O(N^2)$ on worst-case scenarios.
- Heapsort is always $O(N \cdot \log(N))$ but it takes 4-5 times longer to sort typical scenario.

Introsort

- Quicksort is very fast for most scenarios, but it can become $O(N^2)$ on worst-case scenarios.
- Heapsort is always $O(N \cdot \log(N))$ but it takes 4-5 times longer to sort typical scenario.

Introsort does BOTH!

Introsort

Quicksort recurses on $2 * \log(N)$ levels max.

Then Heapsort is called on the rest in case it's still not sorted.

$O(N \cdot \log(N))$ in all cases.

```
template<typename RandIt, typename _Compare>
inline void
sort(RandIt first, RandIt last,
     _Compare comp) {
    if (first != last){
        std::__introsort_loop(first, last,
                               std::__lg(last-first)*2,
                               comp);
        std::__final_insertion_sort(first, last, comp);
    }
}
```

Insertion Sort

$O(N^2)$ algorithm.

Insertion Sort

$O(N^2)$ algorithm.

Over small subranges, Insertion Sort performs better than Quicksort.

We sort until the size of subranges are $< k$.

```
std::nth_element
```

```
template<typename RandIt>
inline void
nth_element(RandIt first, RandIt nth,
            RandIt last)
{
    if (first == last || nth == last)
        return;

    std::__introselct(first, nth, last,
                      std::__lg(last - first) * 2,
                      __gnu_cxx::__ops::__iter_less_iter());
}
```

Introselect

- Quickselect up to $2 * \log(N)$ recursions.
- Then Heapsselect.

Quickselect

Simplified Quicksort.

We choose a pivot and we partition like Quicksort.

Quickselect

Simplified Quicksort.

We choose a pivot and we partition like Quicksort.

If the pivot ends up at n th position, we are done.

Quickselect


Simplified Quicksort.

We choose a pivot and we partition like Quicksort.


If the pivot ends up at n th position, we are done.

Otherwise, we recurse only on the side that contains n th position.

Quickselect



Quickselect



Heapselect

We create a heap on n elements.

Heapselect

We create a heap on n elements.

We iterate on all other elements.

Heapselect

We create a heap on n elements.

We iterate on all other elements.

If an element is smaller than heap max, we pop the heap and add the new element.

Heapselect

We create a heap on n elements.

We iterate on all other elements.

If an element is smaller than heap max, we pop the heap and add the new element.

Heap max is the n th element at the end.

Heapselect

- $O(N \cdot \log(k))$.

```
std::partial_sort
```

GCC Implementation

```
template<typename RandIt, typename _Compare>
inline void
__partial_sort(RandIt first,
               RandIt middle,
               RandIt last,
               _Compare comp)
{
    std::__heap_select(first, middle, last, comp);
    std::__sort_heap(first, middle, comp);
}
```


Partial sort

We said that heapsort was slower than quicksort.

Heapselect is $O(N \cdot \log(k))$ when quickselect is $O(N)$.

How come `partial_sort` can sometimes perform better than quickselect + quicksort?!

Heapselect benchmark

- 1.000.000 elements
- We vary the position we are searching

Heapselect benchmark



Partial sort

- Heapsselect performs better with low k value
- Heapsselect has bigger complexity.

Context of usage

Nth element

Nth element has valid use-cases for any value of k .

First decile, last decile, median...

The STL implementers chose a $O(N)$ algorithm that doesn't depend on k .

Partial Sort

The typical usage of `std::partial_sort` is to sort a small subset of elements in a big container.

The STL implementers chose a faster $O(N \cdot \log(k))$ algorithm that performs well for this typical use-case at the expense of other scenarios.

Despite the fact that the STL is generic, some choices have to be made.

In this case, algorithms are fine-tuned for their typical use cases.

The STL implementers knew what they were doing!

And now we do too!

A Little Order!

Delving into the STL sorting algorithms

Fred Tingaud

@FredTingaudDev