

# RAPID PROTOTYPING OF GRAPHICS SHADERS IN MODERN C++



VALENTIN GALEA  
[@valentin\\_galea](https://twitter.com/valentin_galea)

More than 10 years in *mobile, indie and AAA games*

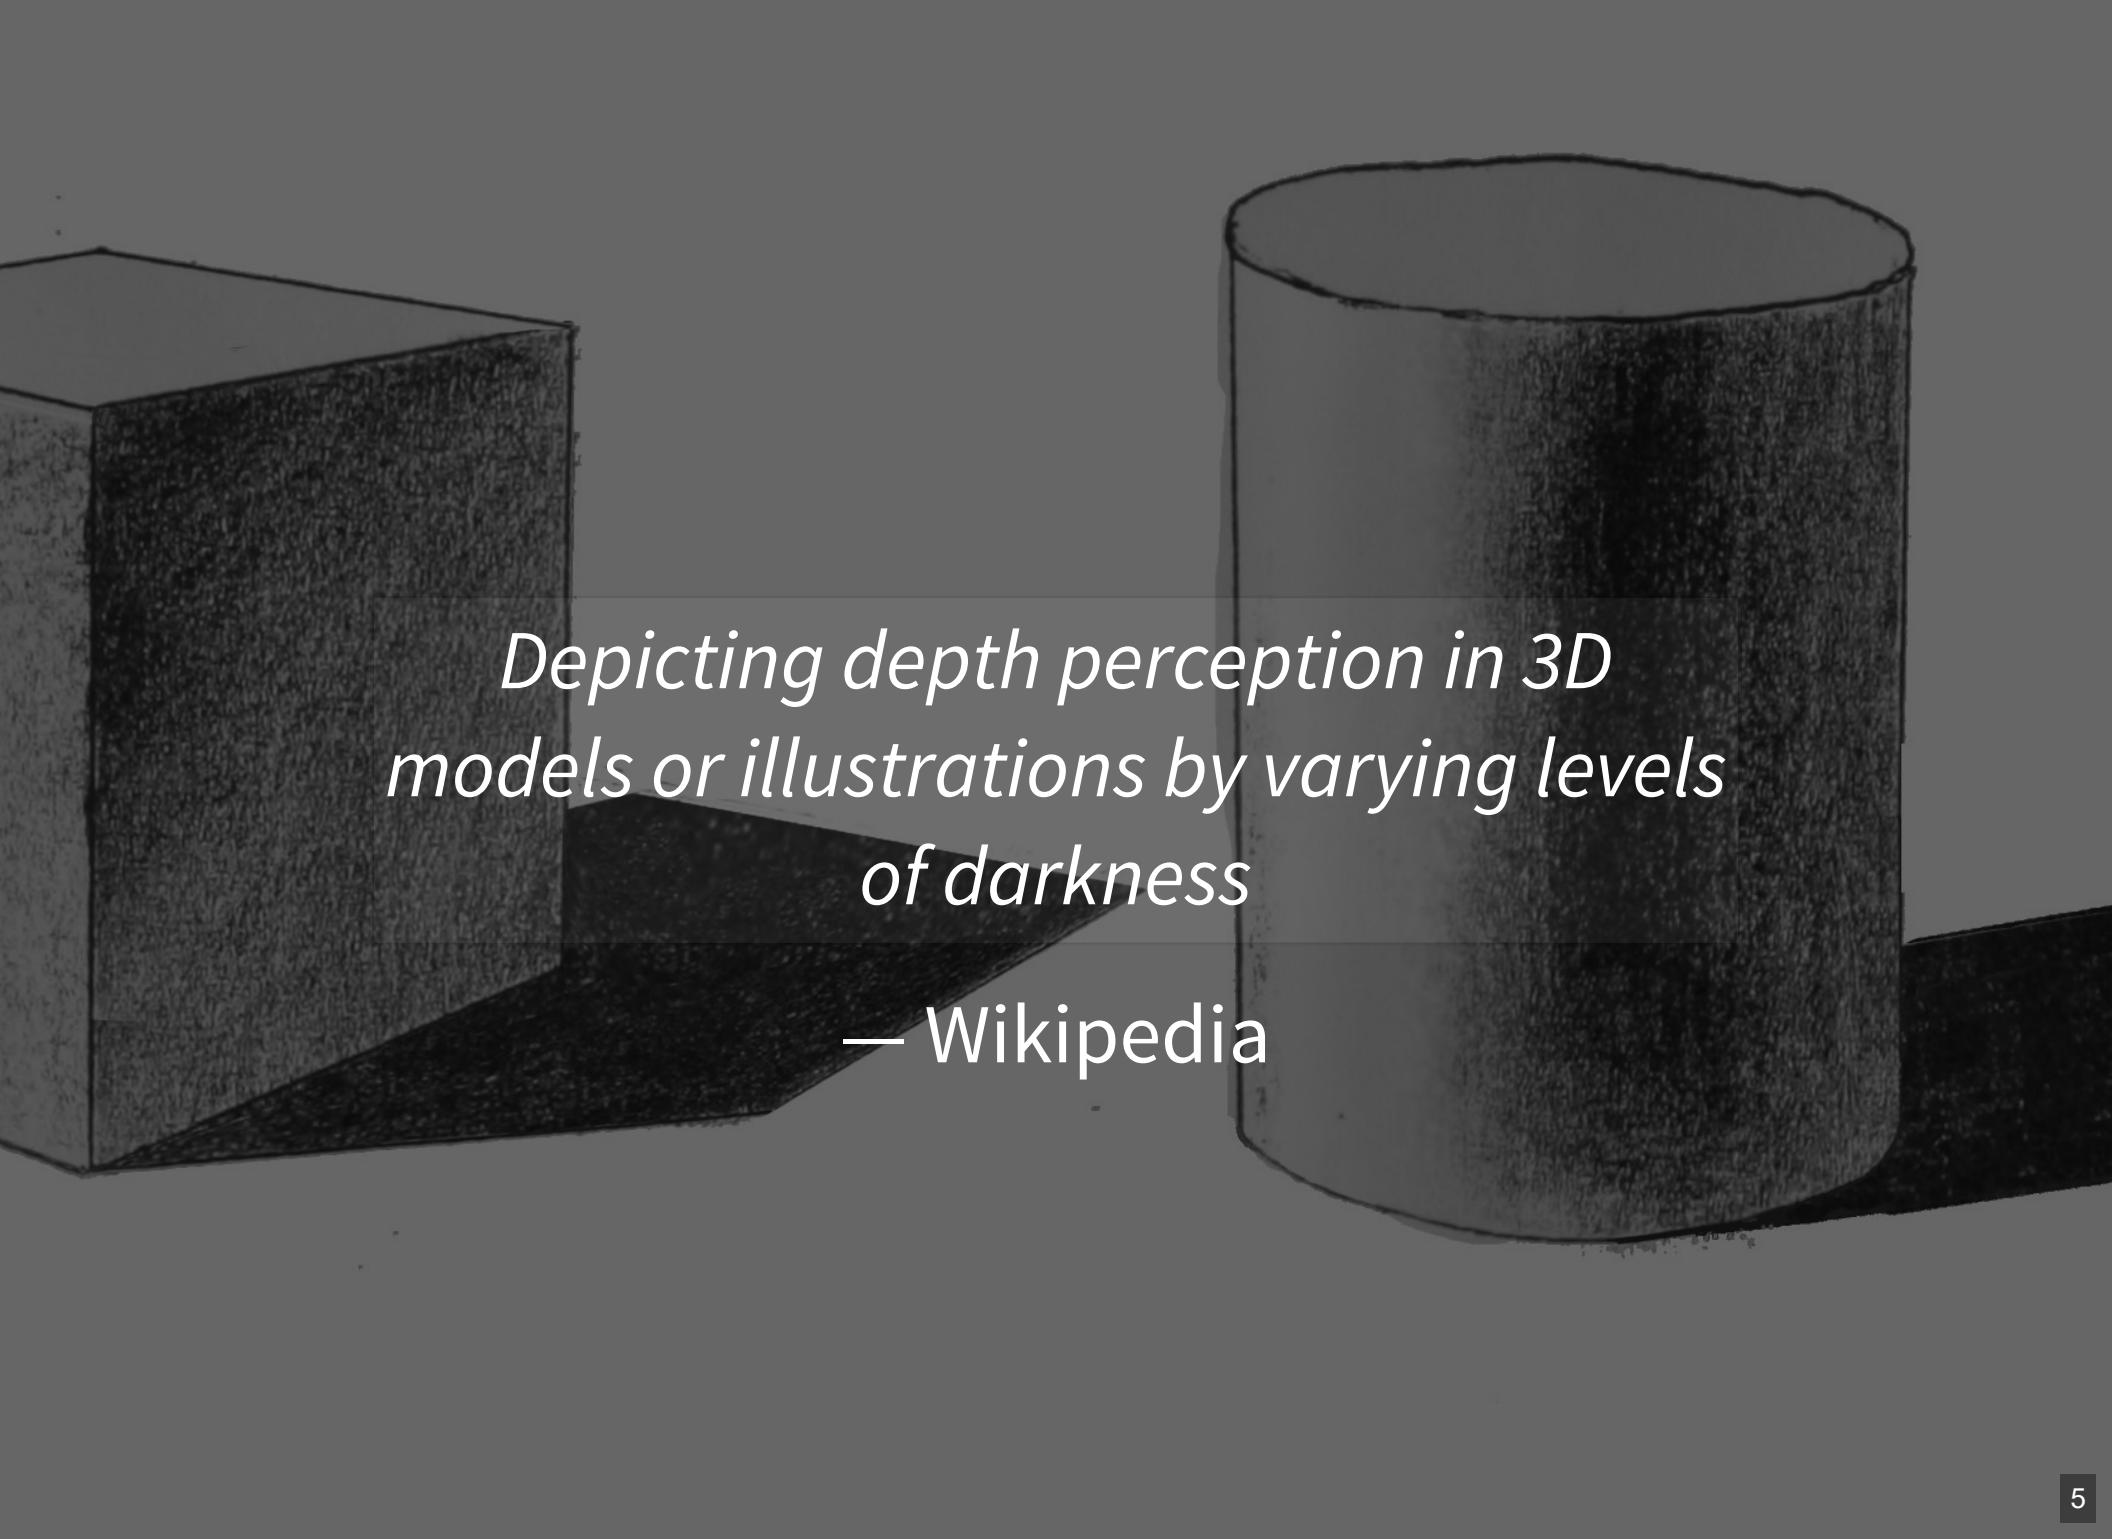


# AGENDA

- Intro and Motivation
- Shading Languages
- C++
- Showcase

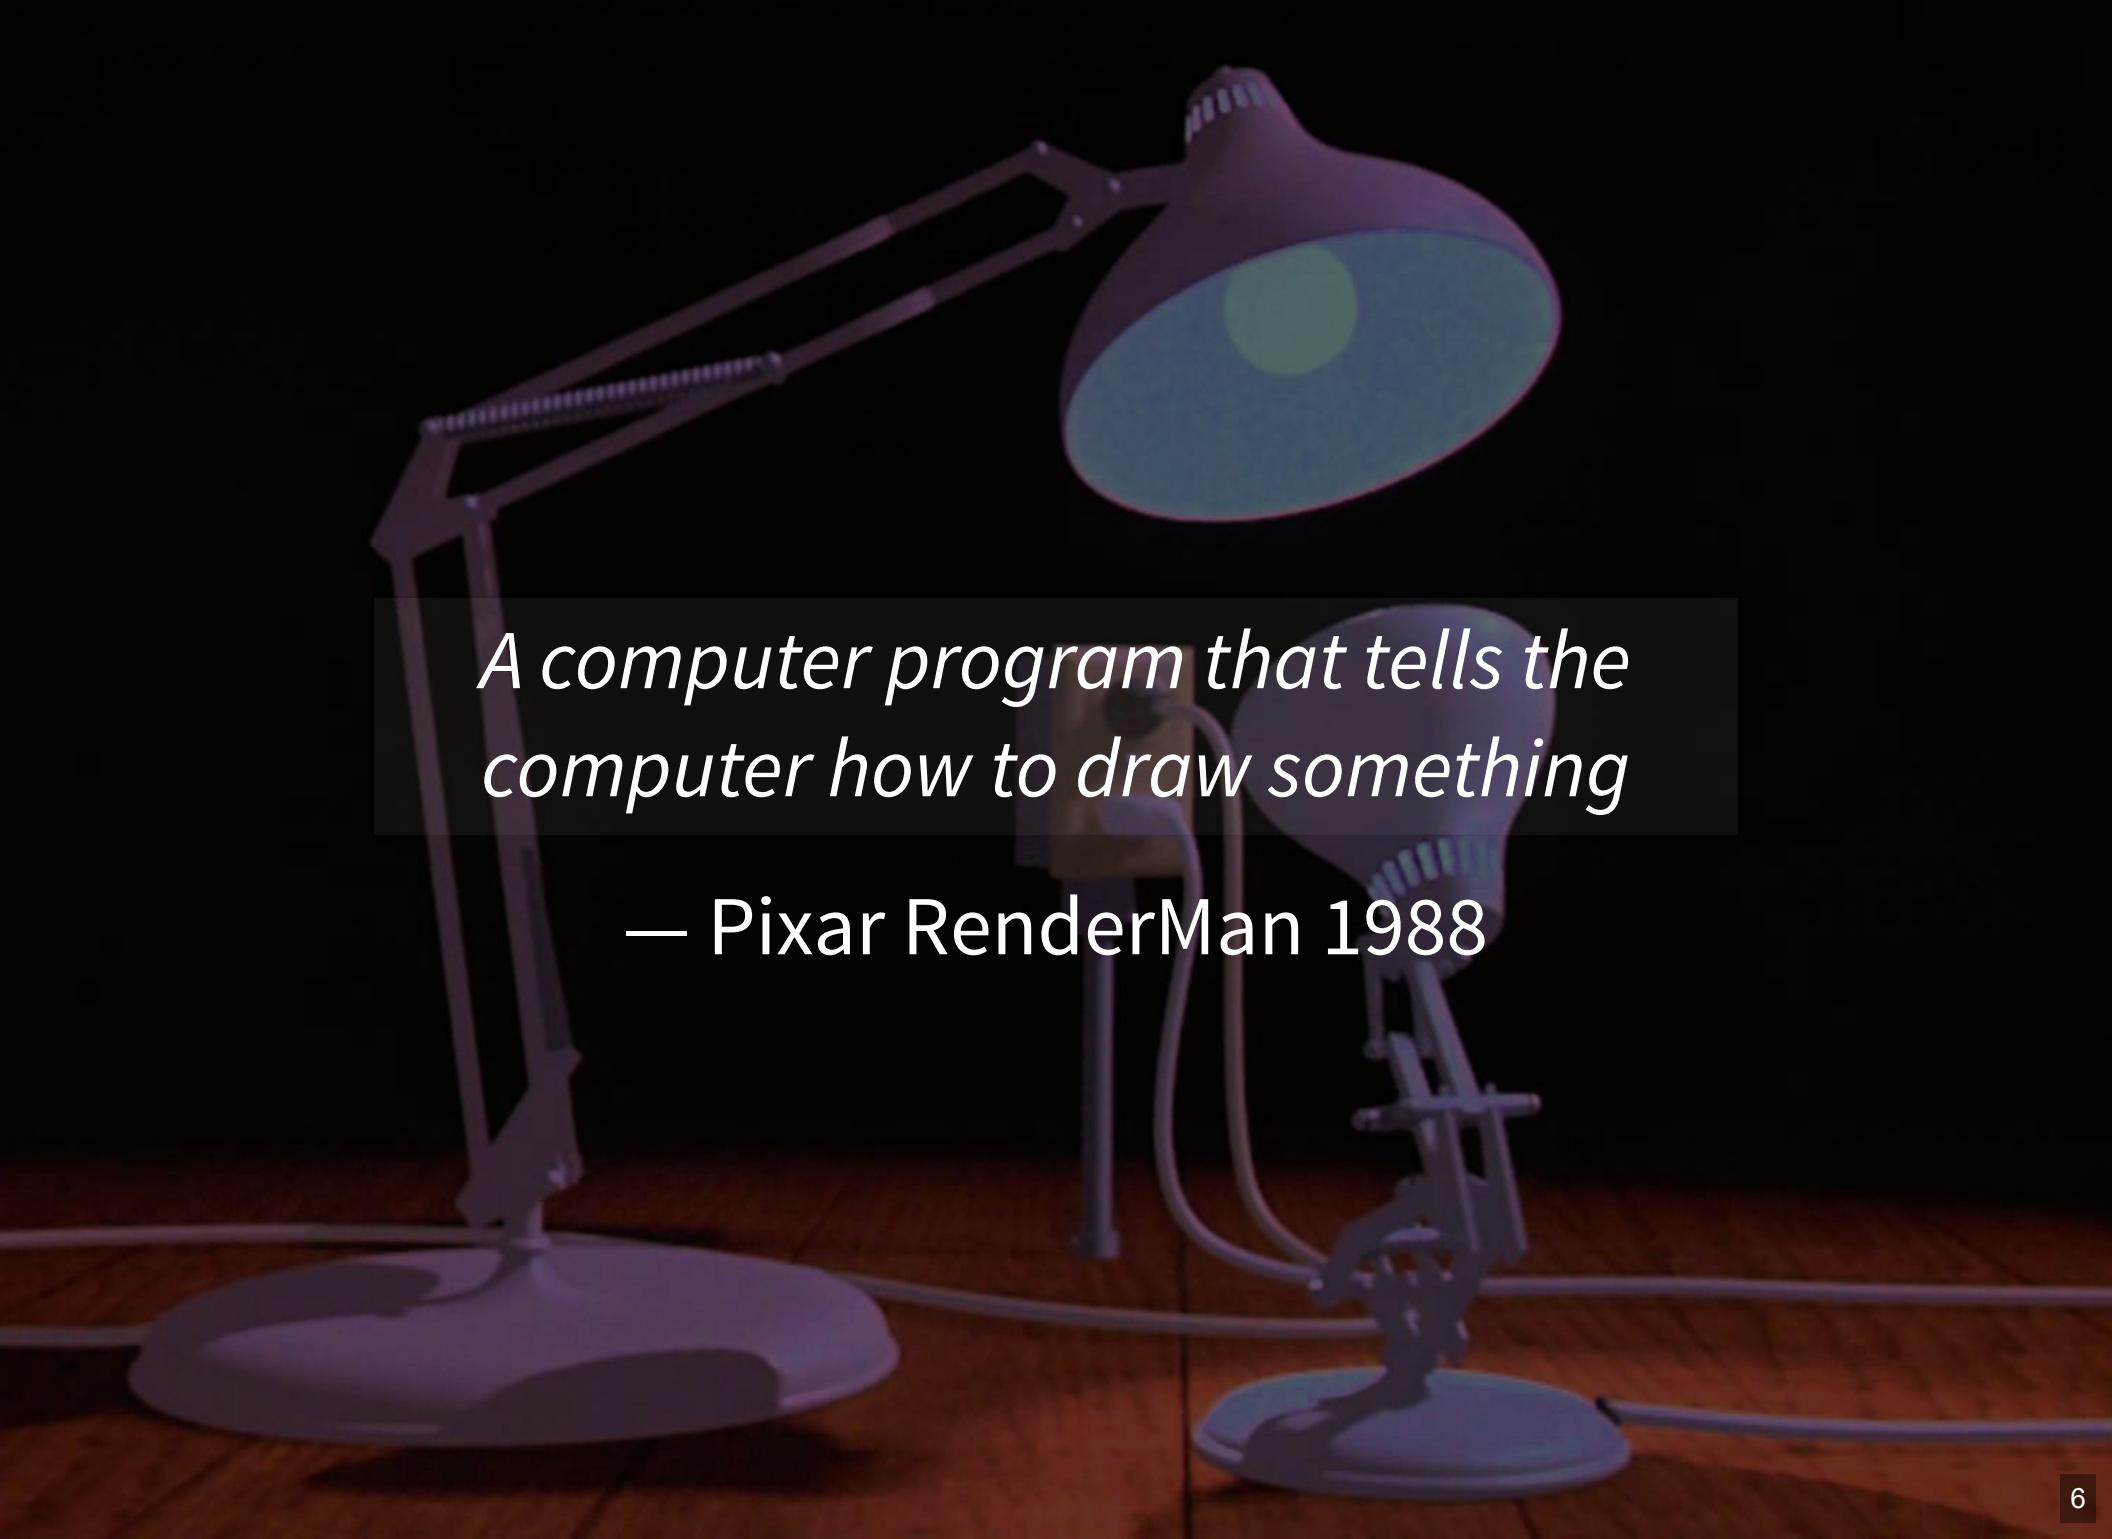
# SHADERS

*A computer program that is used to do shading*

A grayscale photograph showing a dark, textured cylindrical object on the right and a lighter, textured rectangular object on the left. They appear to be resting on a surface.

*Depicting depth perception in 3D  
models or illustrations by varying levels  
of darkness*

– Wikipedia

A Pixar RenderMan 1988 scene set in a dark room with wooden floors. Two lamps are on a desk: a desk lamp with a purple shade and a blue base, and a floor lamp with a green shade and a grey base. A power cord runs from the floor lamp across the floor.

*A computer program that tells the  
computer how to draw something*

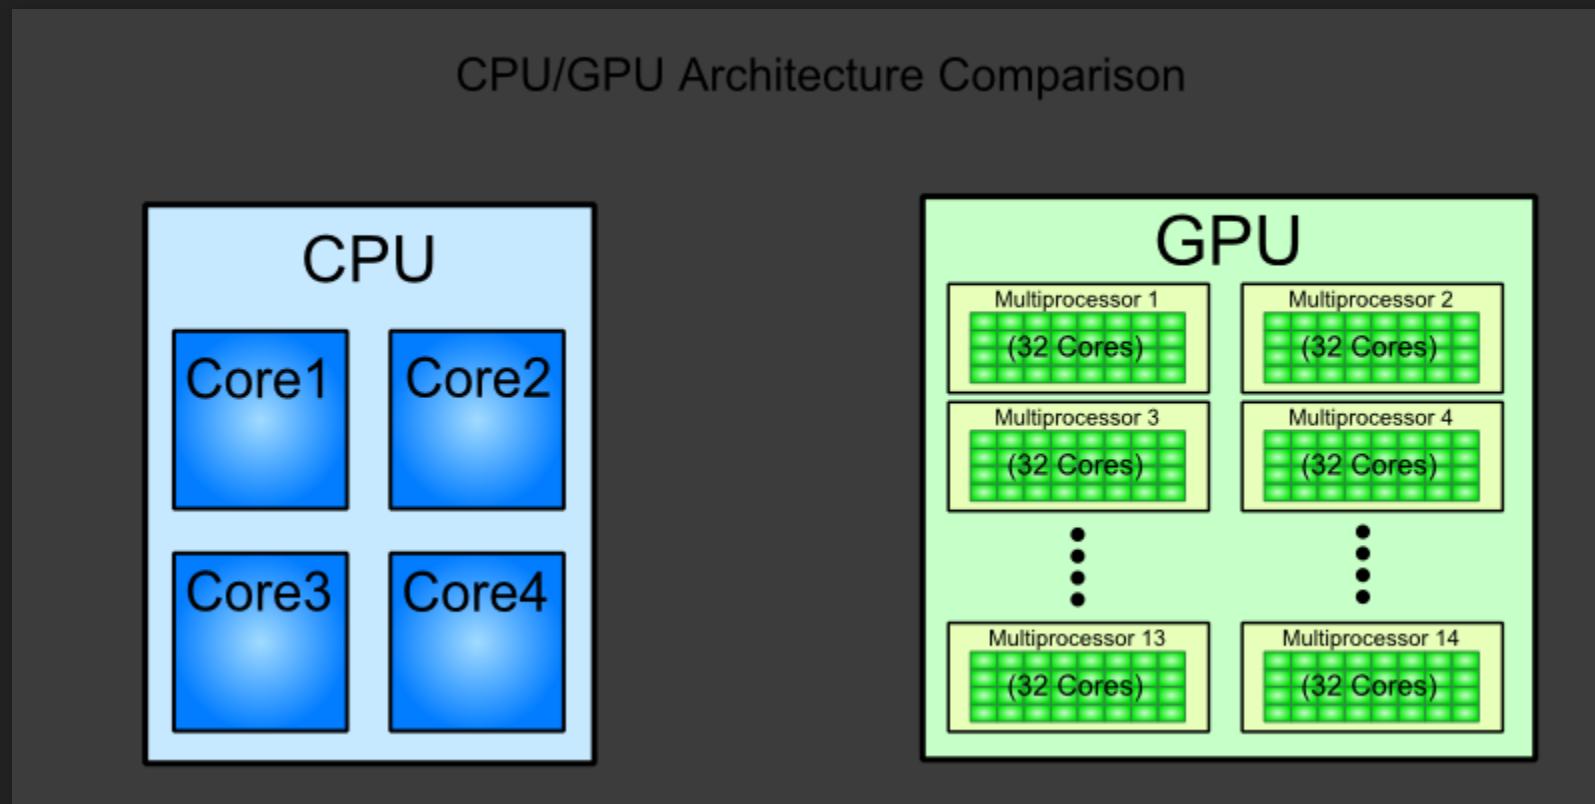
— Pixar RenderMan 1988

## SHADERS ON MODERN GPU'S

- Computer graphics / Images
- Highly parallel computing
- Mining for cryptocurrency :)

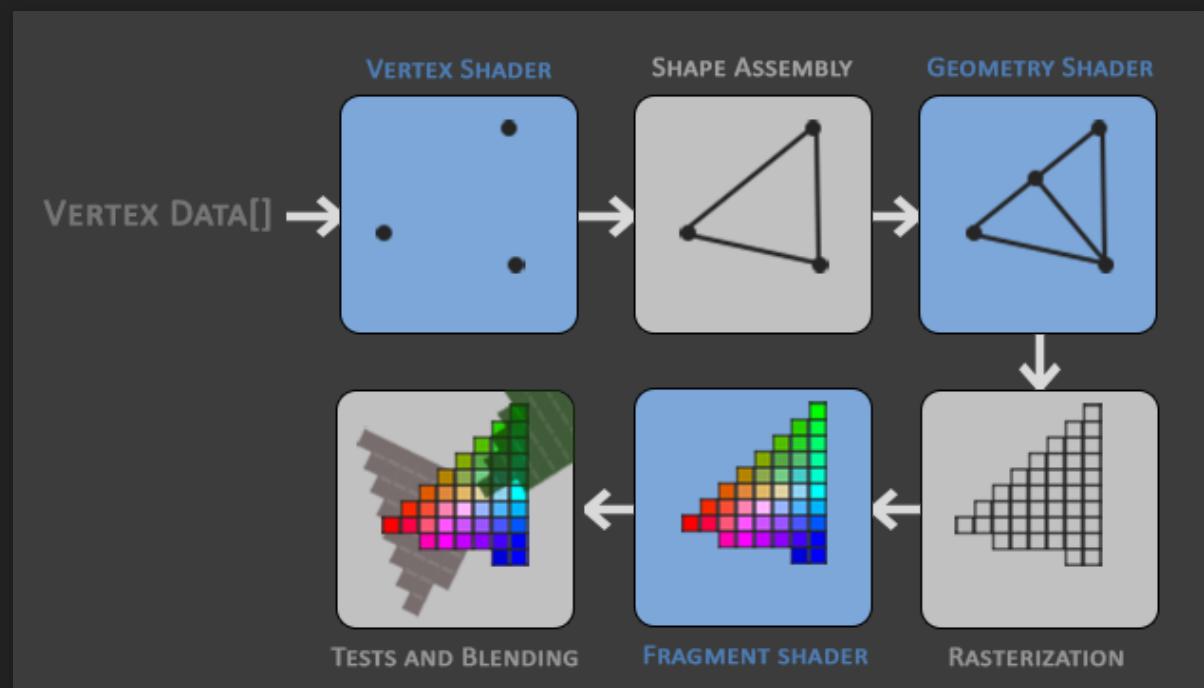


# WHY ON GPU



# TYPES OF SHADERS

Vertex  
Geometry  
Tessellation  
Fragment(Pixel)  
-----  
Compute



# PIXEL/FRAGMENT SHADER



- ! We will concentrate on (procedurally generated) image-only shaders

# MOTIVATION

I wanted to create real-time effects like...

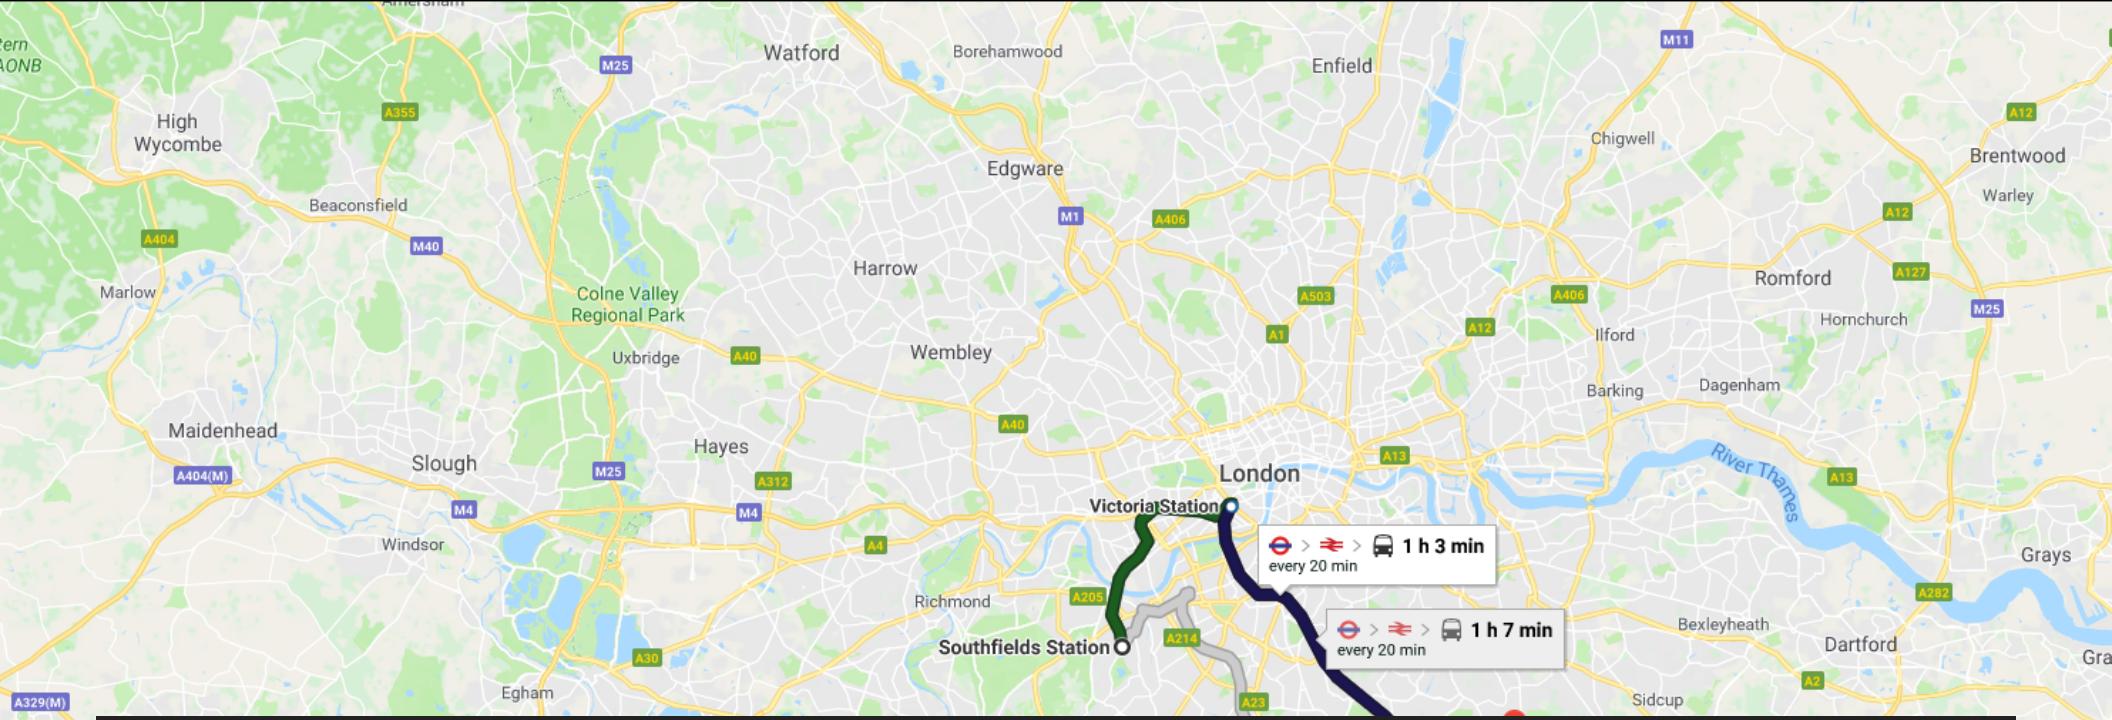




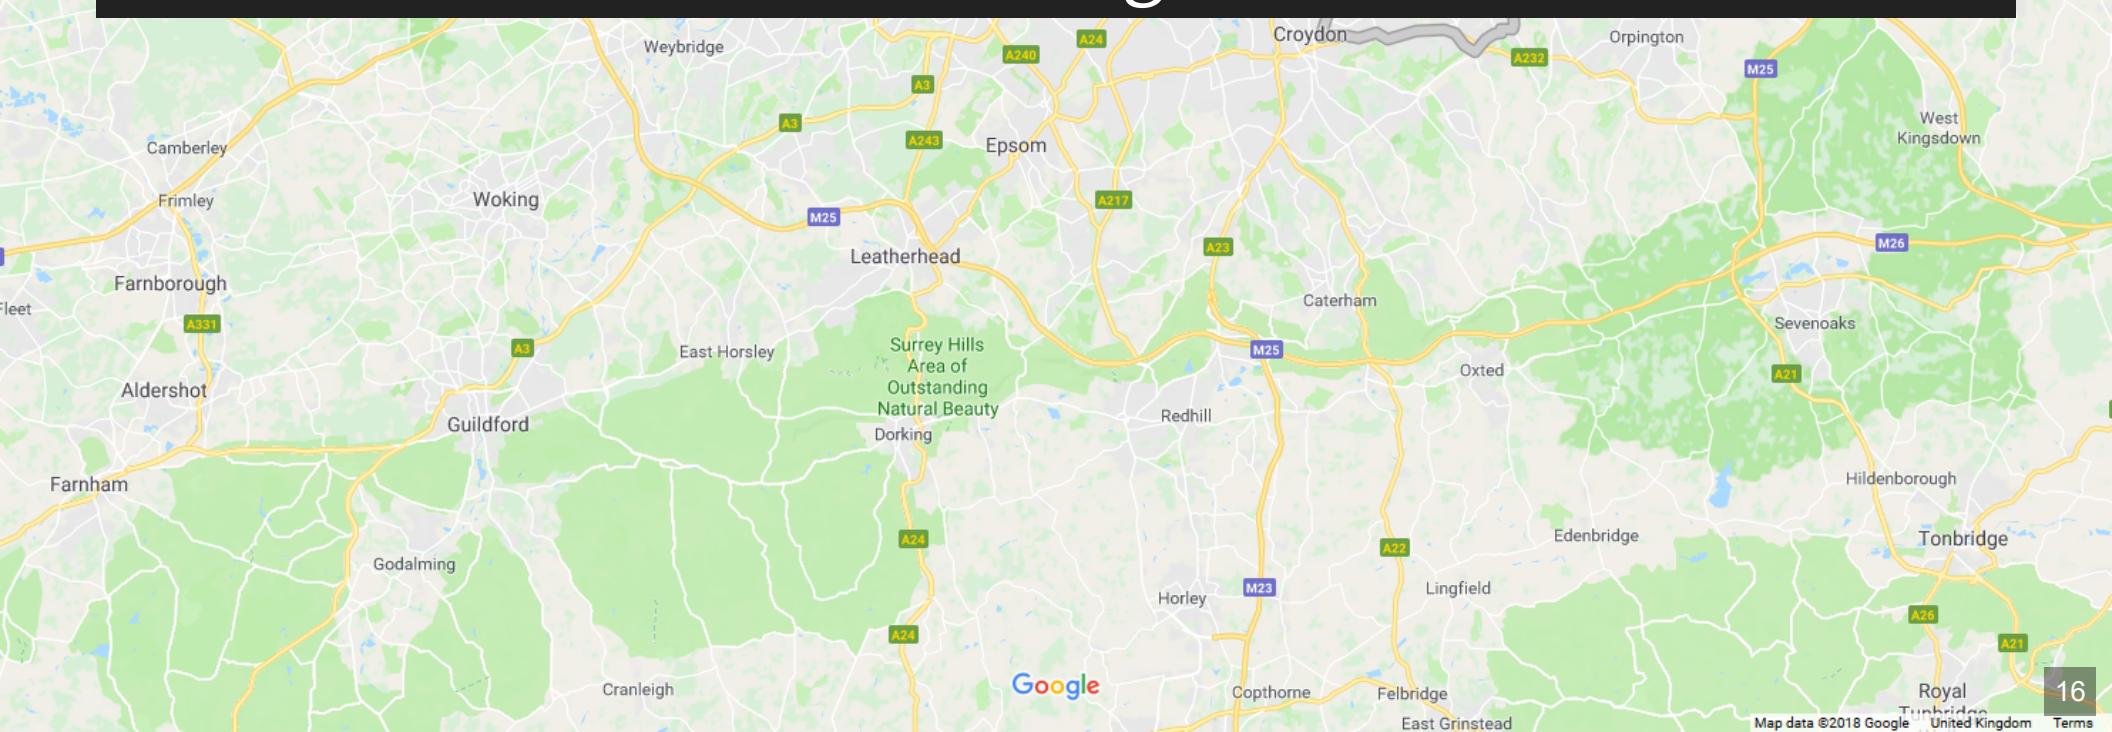


on more limited devices...





...because of long commutes!





also because GPU driver render bugs

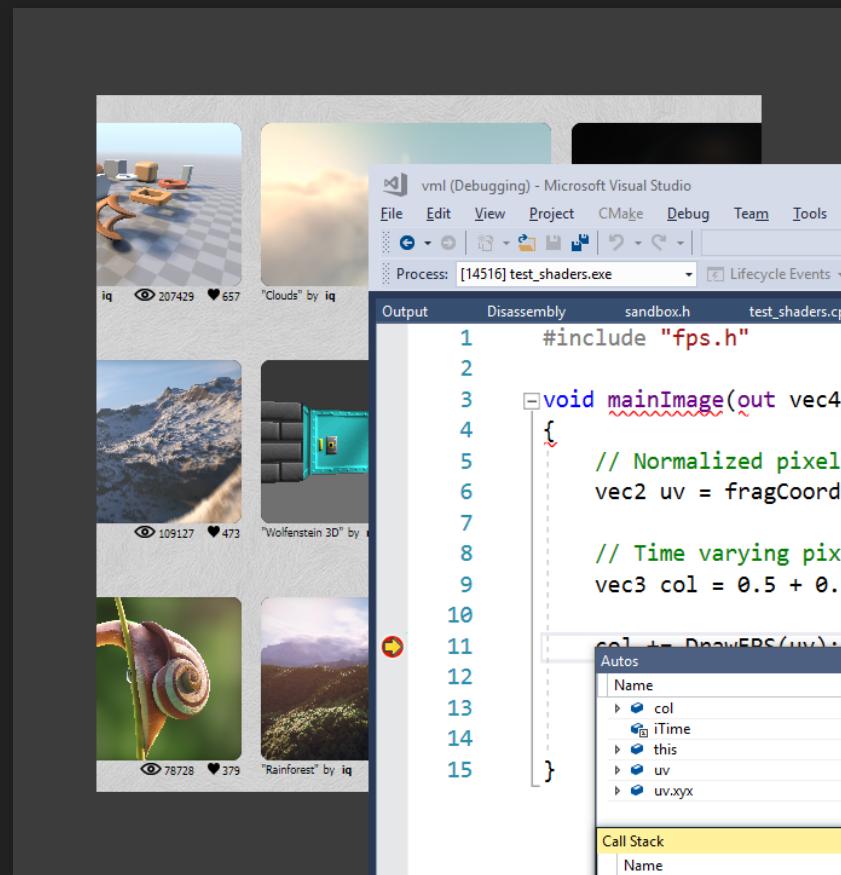


taken on desktop PC / Nvidia GTX 1060

# GIVEAWAY: VML

<https://github.com/valentingalea/vml>

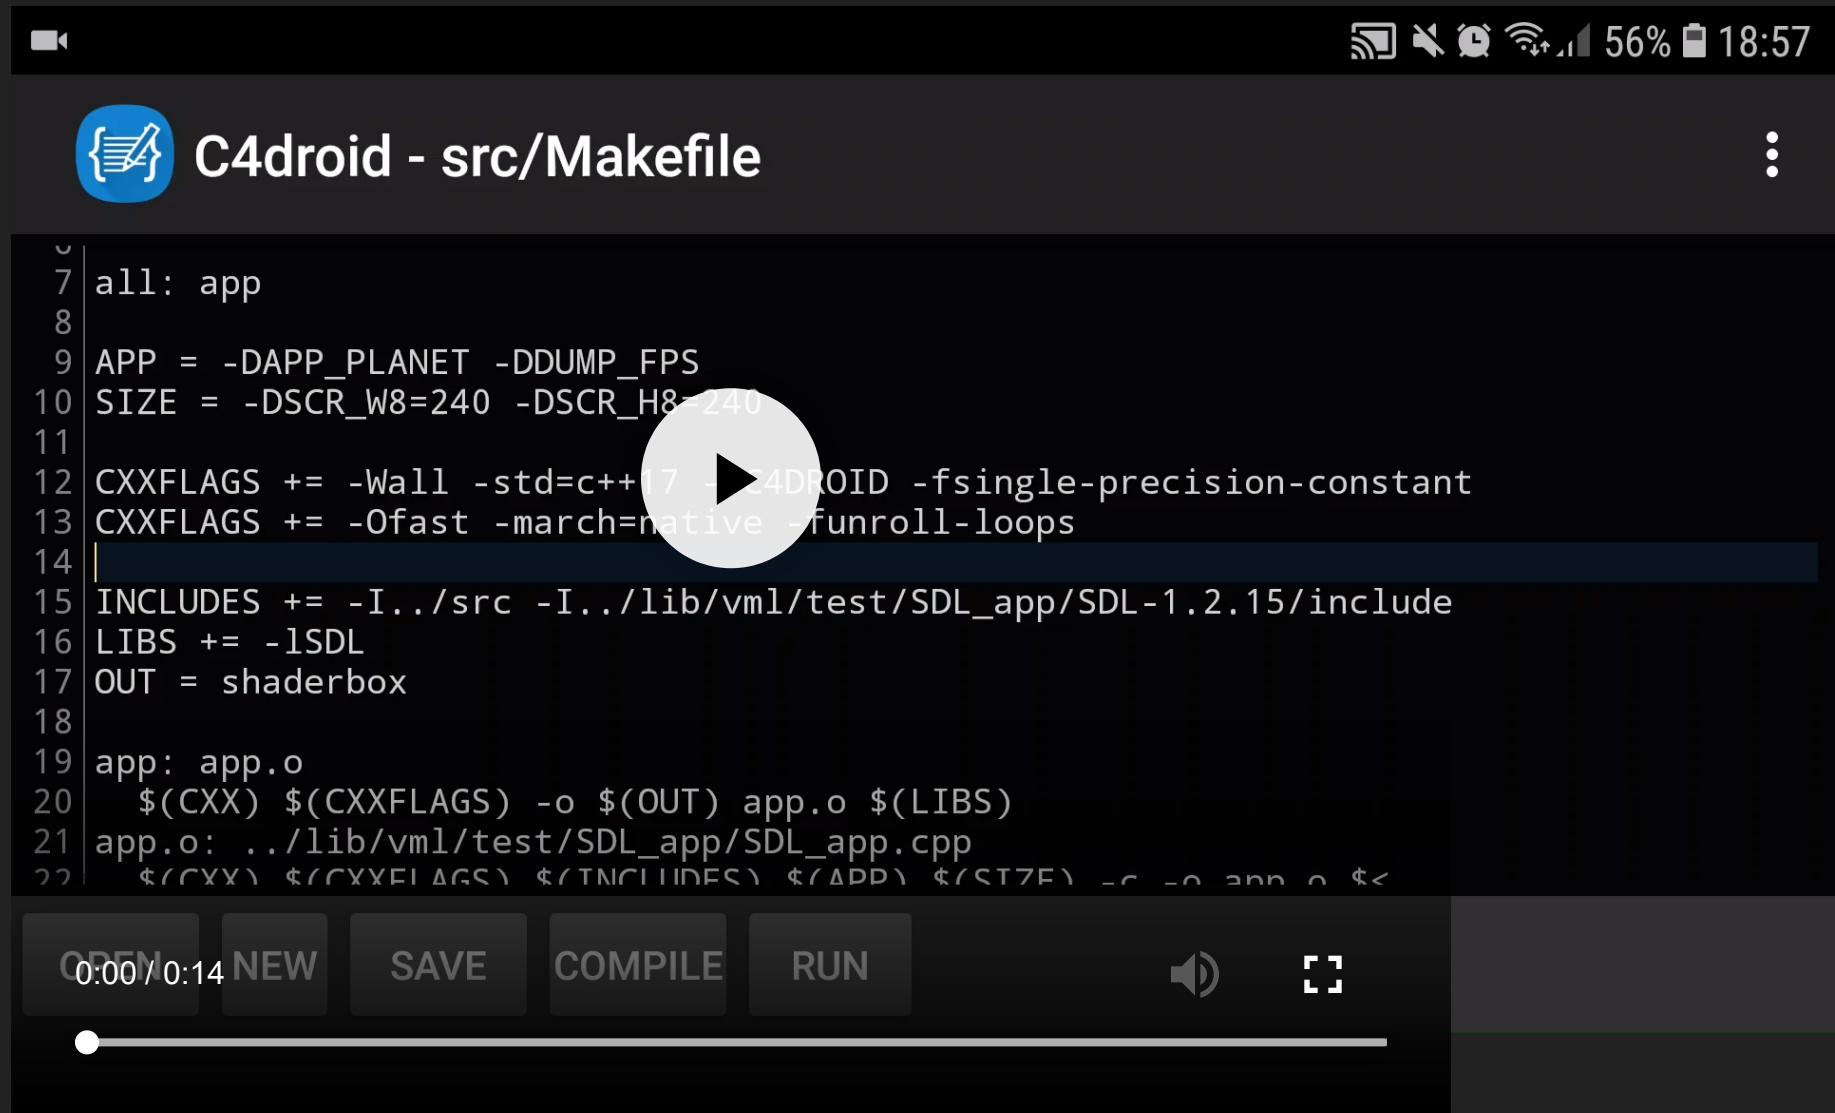
- Debug / Decompile algorithms (from Shadertoy, etc)
- Texture generator (noise patterns, etc)
- Unit Test shaders!
- Quick prototype



# DEMO: GPU VIA SHADERTOY.COM



# DEMO: CPU VIA GCC 8.0 (MOBILE)



The screenshot shows the C4droid application interface on a mobile device. The title bar at the top displays "C4droid - src/Makefile". The main area contains a terminal window showing a Makefile. A large white play button is overlaid on the terminal window, indicating that the code can be run. The Makefile content is as follows:

```
7 all: app
8
9 APP = -DAPP_PLANET -DDUMP_FPS
10 SIZE = -DSCR_W8=240 -DSCR_H8=240
11
12 CXXFLAGS += -Wall -std=c++17 -D4DROID -fsingle-precision-constant
13 CXXFLAGS += -Ofast -march=native -funroll-loops
14
15 INCLUDES += -I../src -I../lib/vml/test/SDL_app/SDL-1.2.15/include
16 LIBS += -lSDL
17 OUT = shaderbox
18
19 app: app.o
20     $(CXX) $(CXXFLAGS) -o $(OUT) app.o $(LIBS)
21 app.o: ../lib/vml/test/SDL_app/SDL_app.cpp
22     $(CXX) $(CXXFLAGS) $(INCLUDES) $(APP) $(ST7E) -c -o app.o $<
```

At the bottom of the screen, there is a toolbar with the following buttons from left to right: OPEN (with a timestamp 0:00 / 0:14), NEW, SAVE, COMPILE, RUN, a volume icon, and a settings icon. A progress bar is also visible at the bottom.

# SHADING LANGUAGES

# PIXAR RENDERMAN LANGUAGE

```
/*
 * red mesh
 */
surface basic() {
    Ci = (1.0, 0.0, 0.0);
    Oi = 1;
}

red shaded mesh

surface simple(color myOpacity = 1) {
    color myColor = (1.0, 0.0, 0.0);
    normal Nn = normalize(N);
    Ci = myColor * myOpacity * diff;
    Oi = myOpacity;
}
```

# SHADING LANGUAGES HISTORY

For real-time rendering:

- Early: ARB assembly, Cg
- OpenGL shading language (**GLSL**)
- DirectX High-Level Shader Language (**HLSL**)
- PlayStation Shader Language (similar to HLSL)

# GLSL VS HLSL

```
varying vec3 N;
varying vec3 v;
void main(void)
{
    vec3 L = normalize(gl_LightSource[0].position.xyz - v);
    vec4 Idiff = gl_FrontLightProduct[0].diffuse * max(dot(N,L), 0);
    Idiff = clamp(Idiff, 0.0, 1.0);
    gl_FragColor = Idiff;
}
```

```
float4 main(
    float3 Light: TEXCOORD0,
    float3 Norm : TEXCOORD1) : COLOR
{
    float4 diffuse = { 1.0, 0.0, 0.0, 1.0 };
    float4 ambient = { 0.1, 0.0, 0.0, 1.0 };
    return ambient + diffuse * saturate(dot(Light, Norm));
}
```

# LANGUAGES: TYPES

GLSL

HLSL

---

Scalar    `bool, int, uint, float, double`

---

Vector    `vec2, vec3, ...`    `float2, float3, ...`

---

Matrix    `mat2, mat3, ...`    `float2x2, float3x3,`

...

---

...textures, samplers, precision modifiers etc

# LANGUAGES: DECLARATIONS

GLSL

HLSL

C++

---

C-style types and arrays

---

C-style struct

---

```
T name = T (   T name = { ... both  
...  } }
```

# LANGUAGES: FUNC ARGS

GLSL and HLSL

C++

---

in T

---

T

out T

---

T &

---

inout T

---

T &

---

const [...] T

# VECTORS AND MATRICES

```
// vectors are generic
vec2 texcoord1, texcoord2;
vec3 position;
vec4 myRGBA;
ivec2 textureLookup;
bvec3 less;
```

```
// matrices are floating point only
mat2 mat2D;
mat3 optMatrix;
mat4 view, projection;
```

# VECTOR SWIZZLE

Syntactic sugar for easy referring to components (or combination of)

{ x, y, z, w }

to represent points or normals

---

{ r, g, b, a }

to refer to colors (a is alpha/translucency)

---

{ s, t, p, q }

texture coordinates

# VECTOR SWIZZLE - EXAMPLES

subcomponents mix & match

```
vec4 v4;  
v4.rgb;    // is a vec4 and the same as just using v4,  
v4.r;     // is a vec3,  
v4.b;     // is a float,  
v4.xy;    // is a vec2,
```

```
vec4 pos = vec4(1.0, 2.0, 3.0, 4.0);  
vec4 swiz= pos.wzyx; // swiz = (4.0, 3.0, 2.0, 1.0)  
vec4 dup = vec4(pos.xx, pos.yy);
```

l-value assignment

```
pos.xw = vec2(5.0, 6.0); // pos = (5.0, 2.0, 3.0, 6.0)  
pos.xx = vec2(3.0, 4.0); // illegal - 'x' used twice
```

# VECTOR SWIZZLE - MOTIVATION

```
vec3 calcNormal( in vec3 pos )
{
    vec2 e = vec2(1.0, -1.0) * 0.0005;

    return normalize(
        e.xyy * map( pos + e.xyy ).x +
        e.yyx * map( pos + e.yyx ).x +
        e.yxy * map( pos + e.yxy ).x +
        e.hxx * map( pos + e.hxx ).x );
}
```

# OPERATORS

syntax

equivalent

```
w = v + u;
```

```
w.x = v.x + u.x;  
w.y = v.y + u.y;  
w.z = v.z + u.z;
```

```
u = v * m;
```

```
u.x = dot(v, m[0]);  
u.y = dot(v, m[1]);  
u.z = dot(v, m[2]);  
/* dot(a,b) is the inner product of a and b */
```

## "STANDARD LIBRARY"

Math      `sin, cos, radians, pow, exp, etc`

---

Common    `abs, sign, floor, mod, min, etc`

---

Utility    `mix, step, smoothstep, etc`

---

Geometry    `length, dot, cross, distance, etc`

---

Specific texture and image sampling ...

# SHADING LANGUAGES FUTURE

Basically C++ (usually via LLVM)

# SHADING LANGUAGES FUTURE

Basically C++ (usually via LLVM)

- Metal Shading Language (C++14, Apple)
  - only on iOS devices

# SHADING LANGUAGES FUTURE

Basically C++ (usually via LLVM)

- Metal Shading Language (C++14, Apple)
  - only on iOS devices
- CUDA Heterogeneous Computing (C++11, NVidia)
  - only for computing, not graphics

# SHADING LANGUAGES FUTURE

Basically C++ (usually via LLVM)

- Metal Shading Language (C++14, Apple)
  - only on iOS devices
- CUDA Heterogeneous Computing (C++11, NVidia)
  - only for computing, not graphics
- HLSL 6.x (C++98'ish, Microsoft)
  - not released yet

Let's see how C++ can help out, NOW!

# THE PLAN

## THE PLAN

-  Pick a shading language and twist C++ to accept it as source code!

## THE PLAN

-  Pick a shading language and twist C++ to accept it as source code!
-  BONUS: use the preprocessor for transcription back to the original language(s)!

# THE PLAN (CONT.)

## THE PLAN (CONT.)

- obligatory preprocessor layer

## THE PLAN (CONT.)

- obligatory preprocessor layer
- vector (linear algebra) types
  - swizzle support

## THE PLAN (CONT.)

- obligatory preprocessor layer
- vector (linear algebra) types
  - swizzle support
- matrix types

## THE PLAN (CONT.)

- obligatory preprocessor layer
- vector (linear algebra) types
  - swizzle support
- matrix types
- operators

## THE PLAN (CONT.)

- obligatory preprocessor layer
- vector (linear algebra) types
  - swizzle support
- matrix types
- operators
- "standard library" utility/math functions

## PLACE YOUR BETS

We will chose **GSL** as it's used on *desktop, web and mobile*

- ⚠ only a subset of it - concentrate on procedural graphics thus minimize/eliminate inputs (textures, vertex data, etc)

# DESIGN OF vector<>

# DESIGN OF `vector<>`

```
template<typename T, size_t N>
struct vector :
    public vector_base<T, N>
{
```

# DESIGN OF `vector<>`

```
template<typename T, size_t N>
struct vector :
    public vector_base<T, N>
{
    vector();
    explicit vector(scalar_type s);
    template<typename... Args> explicit vector(Args... args);
```

# DESIGN OF `vector<>`

```
template<typename T, size_t N>
struct vector :
    public vector_base<T, N>
{
    vector();
    explicit vector(scalar_type s);
    template<typename... Args> explicit vector(Args... args);

    scalar_type& operator[](size_t i);
    scalar_type& operator[](size_t i);

    vector_type& operator+=(scalar_type s);
    vector_type& operator+=(const vector_type& v);
(etc)
```

# vector<> CTOR - BASIC

```
vector()
{
    static_for<0, N>() ([this] (size_t i) {
        data[i] = 0;
    }) ;
}
```

```
explicit vector(scalar_type s)
{
    static_for<0, N>() ([s, this] (size_t i) {
        data[i] = s;
    }) ;
}
```

**static\_for UTILITY**

# static\_for UTILITY

```
template<size_t Begin, size_t End>
struct static_for
{
    template<class Func>
    constexpr void operator ()(Func&& f)
    {
        f(Begin);
        static_for<Begin + 1, End>() (std::forward<Func>(f));
    }
};
```

# static\_for UTILITY

```
template<size_t Begin, size_t End>
struct static_for
{
    template<class Func>
    constexpr void operator ()(Func&& f)
    {
        f(Begin);
        static_for<Begin + 1, End>() (std::forward<Func>(f));
    }
};
```

```
template<size_t N>
struct static_for<N, N>
{
    template<class Func>
    constexpr void operator ()(Func&&) { /* terminate */ }
};
```

# **vector<> CTOR - ADVANCED**

# vector<> CTOR - ADVANCED

```
template<typename A0, typename... Args,
```

## vector<> CTOR - ADVANCED

```
template<typename A0, typename... Args,  
         class = typename std::enable_if<  
             ((sizeof... (Args) >= 1) ||  
              ((sizeof... (Args) == 0) && !std::is_scalar_v<A0>))  
         >::type>
```

# vector<> CTOR - ADVANCED

```
template<typename A0, typename... Args,  
  
         class = typename std::enable_if<  
             ((sizeof... (Args) >= 1) ||  
              ((sizeof... (Args) == 0) && !std::is_scalar_v<A0>))  
         >::type>  
  
explicit vector(A0&& a0, Args&&... args)  
{  
    static_recurse<0>(  
        std::forward<A0>(a0),  
        std::forward<Args>(args)...  
    );  
}
```

## **vector<> CTOR - ADVANCED (2)**

## vector<> CTOR - ADVANCED (2)

```
template<size_t I, typename Arg0, typename... Args>
void static_recurse(Arg0&& a0, Args&&... args)
{
    construct_at_index<I>(std::forward<Arg0>(a0)) ;

    static_recurse<I + get_size<Arg0>()>(
        std::forward<Args>(args)...
    );
}
```

## vector<> CTOR - ADVANCED (2)

```
template<size_t I, typename Arg0, typename... Args>
void static_recurse(Arg0&& a0, Args&&... args)
{
    construct_at_index<I>(std::forward<Arg0>(a0)) ;

    static_recurse<I + get_size<Arg0>()>(
        std::forward<Args>(args)...
    );
}
```

```
template<size_t I>
void static_recurse()
{ /* terminate */ }
```

## **vector<> CTOR - ADVANCED (3)**

## vector<> CTOR - ADVANCED (3)

```
template<size_t i>
void construct_at_index(scalar_type arg)
{
    data[i] = arg;
}
```

## vector<> CTOR - ADVANCED (3)

```
template<size_t i>
void construct_at_index(scalar_type arg)
{
    data[i] = arg;
}
```

```
template<size_t i, typename Other, size_t Other_N>
void construct_at_index(vector<Other, Other_N>&& arg)
{
    constexpr auto count = std::min(i + Other_N, num_components);

    static_for<i, count>()([&](size_t j) {
        data[j] = arg.data[j - i];
    });
}
```

# vector<> CTOR IN ACTION

```
using vec2 = vector<int, 2>;
using vec3 = vector<int, 3>

vec3 v = vec3(98, vec2(99, 100));
//          ^      ^
//          |      |
//          `-- scalar construct gets called
//          |
//          `---- sub-vector construct gets called
//                  and then recursively again
```

# Godbolt

```
int main()
{
    float a, b;
    scanf("%f %f", &a, &b);

    auto v = vec3(1.f, vec2(a, b));

    printf("%f %f", v.x, v.y);
}
```

-std=c++17 -Wall -O2 (source)

# GODBOLT (CONT.)

5.x/6.x)

gcc (7.x/8.x)

msvc (2017)

```
if
    dword ptr [rsp + 4]
        xmm1, xmm0
    dword ptr [rsp]
        xmm2, xmm0
    qword ptr [rip + .LCPI0_0]
    offset .L.str.1
3
ntf
```

```
call    scanf
pxor   xmm2, xmm2
pxor   xmm1, xmm1
movsd  xmm0, QWORD PTR .LC1[rip]
mov    edi, OFFSET FLAT:.LC2
mov    eax, 3
cvtss2sd      xmm2, DWORD PTR [rsp+12]
cvtss2sd      xmm1, DWORD PTR [rsp+8]
call    printf
```

```
call    scanf
movss  xmm1, DWORD PTR b$`r14
lea    rcx, OFFSET FLAT:`r15
movss  xmm0, DWORD PTR a$`r16
movss  DWORD PTR $T1[rsp+12]
movsd  xmm1, QWORD PTR _r17
movss  DWORD PTR $T1[rsp+8]
movq   rdx, xmm1
mov    rax, QWORD PTR $T1[rsp+16]
mov    QWORD PTR <args_1>[rcx]
movss  xmm3, DWORD PTR <a$`r18>
movss  xmm2, DWORD PTR <a$`r19>
cvtps2pd xmm3, xmm3
cvtps2pd xmm2, xmm2
movq   r9, xmm3
movq   r8, xmm2
call    printf
```

# PROBLEM: DEBUG

```
void static_for<0ul, 3ul>::operator()<vector<float, 3ul>...
push    rbp
mov     rbp, rsp
sub    rsp, 32
mov     QWORD PTR [rbp-24], rdi
mov     QWORD PTR [rbp-32], rsi
mov     rax, QWORD PTR [rbp-32]
mov     esi, 0
mov     rdi, rax
call    vector<float, 3ul>::vector(float)::{lambda(unsigned...
mov     rax, QWORD PTR [rbp-32]
mov     rdi, rax
call    vector<float, 3ul>::vector(float)::{lambda(unsigned...
mov     rdx, rax
lea     rax, [rbp-1]
mov     rsi, rdx
mov     rdi, rax
call    void static_for<1ul, 3ul>::operator()<vector...
nop
leave
ret
```

# SOLUTION: C++17 FOLD EXPRESSIONS

Unary right fold

$(E \text{ op } ...) \rightarrow (E_1 \text{ op } (\dots \text{ op } (E_{N-1} \text{ op } E_N)))$

---

clang 3.6+

gcc 6+

MSVC 2017 15.5+

# IMPROVED DESIGN OF `vector<>`

# IMPROVED DESIGN OF `vector<>`

```
template<typename T, size_t... Ns>
struct vector :
    public vector_base<T, Ns...>
{
```

# IMPROVED DESIGN OF `vector<>`

```
template<typename T, size_t... Ns>
struct vector :
    public vector_base<T, Ns...>
{
    vector()
    {
        ((data[Ns] = 0), ...);
    }
}
```

# IMPROVED DESIGN OF `vector<>`

```
template<typename T, size_t... Ns>
struct vector :
    public vector_base<T, Ns...>
{
```

```
    vector()
    {
        ((data[Ns] = 0), ...);
    }
```

```
    explicit vector(scalar_type s)
    {
        ((data[Ns] = s), ...);
    }
```

```
    template<typename A0, typename... Args> explicit vector
(etc)
```

# FOLDING EXPRESSIONS IN ACTION

declaration

```
explicit vector(scalar_type s)
{
    ((data[Ns] = s), ...);
}
```

instantiation

```
vector<float, 0, 1, 2>

explicit vector(float s)
{
    data[0] = s,
    data[1] = s,
    data[2] = s;
}
```

# IMPROVED vector<> CTOR

```
template<typename A0, typename... Args>
explicit vector(A0&& a0, Args&&... args)
{
    size_t i = 0; // advances as we consume args

    // consume the first one
    construct_at_index(i, std::forward<A0>(a0));

    // consume the rest, if any
    (construct_at_index(i, std::forward<Args>(args)), ...);
}
```

# SWIZZLING

## vector\_base NAIVE IMPL

```
template<typename T>
struct vector_base<T, 2>
{
    union
    {
        T data[2];
        struct { T x, y; };
        struct { T s, t; };
        struct { T u, v; };
    }
};
```

## vector\_base NAIVE IMPL (2)

```
template<typename T>
struct vector_base<T, 3>
{
    union
    {
        T data[3];
        struct { T x, y, z; };
        struct { T r, g, b; };
        struct { T s, t, p; };
    }
};
```

## vector\_base NAIVE IMPL (3)

```
template<typename T>
struct vector_base<T, 4>
{
    union
    {
        T data[4];
        struct { T x, y, z, w; };
        struct { T r, g, b, a; };
        struct { T s, t, p, q; };
    }
};
```

## **vector\_base** NOTES

-  both anonymous struct and union are permitted, only MSVC complains with warning
-  union active member switching can be tricky [10.5] but we'll use only trivial types with trivial assignment

# SWIZZLE

- 💡 We introduce an additional proxy class that allows custom access to the indices and we create all possible permutations (per GLSL/HLSL standard)

```
template<class vector_type, class T, size_t N, size_t... indices>
struct swizzler
{
    T data[N];
    (etc)
```

# SWIZZLE FOR `vector<T, 3>`

```
union
{
    T data[3];

    struct {
        swizzler<0>::type x;
        swizzler<1>::type y;
        swizzler<2>::type z;
    };
    struct {
        swizzler<0>::type r;
        swizzler<1>::type g;
        swizzler<2>::type b;
    };
    struct {
        swizzler<0>::type s;
        swizzler<1>::type t;
        swizzler<2>::type p;
    };
    ...
}
```

## SWIZZLE (CONT.)

```
...
swizzler<0, 0>::type xx, rr, ss;
swizzler<0, 1>::type xy, rg, st;
swizzler<0, 2>::type xz, rb, sp;
swizzler<1, 0>::type yx, gr, ts;
swizzler<1, 1>::type yy, gg, tt;
swizzler<1, 2>::type yz, gb, tp;
swizzler<2, 0>::type zx, br, ps;
swizzler<2, 1>::type zy, bg, pt;
swizzler<2, 2>::type zz, bb, pp;
...
```

## ...MORE SWIZZLE

```
...
swizzler<0, 0, 0>::type xxx, rrr, sss;
swizzler<0, 0, 1>::type xxy, rrg, sst;
swizzler<0, 0, 2>::type xxz, rrb, ssp;
swizzler<0, 1, 0>::type xyx, rgr, sts;
swizzler<0, 1, 1>::type xyy, rgg, stt;
swizzler<0, 1, 2>::type xyz, rgb, stp;
swizzler<0, 2, 0>::type xzx, rbr, sps;
swizzler<0, 2, 1>::type xzy, rbg, spt;
swizzler<0, 2, 2>::type xzz, rbb, spp;
swizzler<1, 0, 0>::type yxx, grr, tss;
swizzler<1, 0, 1>::type yxy, grg, tst;
swizzler<1, 0, 2>::type yxz, grb, tsp;
...
...
```

## ...EVEN MORE SWIZZLE!

```
...
swizzler<2, 1, 2, 0>::type zyzx, bgbr, ptps;
swizzler<2, 1, 2, 1>::type zzyz, bgbg, ptpt;
swizzler<2, 1, 2, 2>::type zyzz, bgbb, ptpp;
swizzler<2, 2, 0, 0>::type zxxx, bbrr, ppss;
swizzler<2, 2, 0, 1>::type zzxy, bbrg, ppst;
swizzler<2, 2, 0, 2>::type zxxz, bbrb, ppsp;
swizzler<2, 2, 1, 0>::type zzyx, bbgr, ppt;
swizzler<2, 2, 1, 1>::type zzyy, bbgg, pptt;
swizzler<2, 2, 1, 2>::type zzyz, bbgb, pptp;
swizzler<2, 2, 2, 0>::type zzzx, bbbr, ppps;
swizzler<2, 2, 2, 1>::type zzyy, bbbg, pppt;
swizzler<2, 2, 2, 2>::type zzzz, bbbb, pppp;
};

}
```

# Swizzler<> DESIGN

```
template<
    typename vector_type,
    typename scalar_type,
    size_t N,
    size_t... indices>
struct swizzler
{
    T data[N];
    // N might differ from vector_type::num_components
    // ex: .xxxx from vec2

(etc)
```

# swizzler<> CONVERSIONS

## swizzler<> CONVERSIONS

Needs to implicitly convert/assign to its `vector<>` equivalent

# swizzler<> CONVERSIONS

Needs to implicitly convert/assign to its `vector<>` equivalent

```
operator vector_type()
{
    vector_type vec;
    assign_across(vec, 0, indices...);
    return vec;
}
```

# swizzler<> CONVERSIONS

Needs to implicitly convert/assign to its `vector<>` equivalent

```
operator vector_type()
{
    vector_type vec;
    assign_across(vec, 0, indices...);
    return vec;
}
```

```
swizzler& operator=(const vector_type& vec)
{
    assign_across(vec, 0, indices...);
    return *this;
}
```

## **swizzler<> DESIGN (CONT.)**

We use same fold expression trick

## swizzler<> DESIGN (CONT.)

We use same fold expression trick

```
template<typename... Indices>
void assign_across(vector_type& vec, size_t i, Indices ...j) const
{
    ((vec[i++] = data[j]), ...);
}
```

## swizzler<> DESIGN (CONT.)

We use same fold expression trick

```
template<typename... Indices>
void assign_across(vector_type& vec, size_t i, Indices ...j) const
{
    ((vec[i++] = data[j]), ...);
}
```

```
template<typename... Indices>
void assign_across(const vector_type& vec, size_t i, Indices ...j)
{
    ((data[j] = vec[i++]), ...);
}
```

# swizzler<> PROBLEM

```
vec3 v = vec4(other.xy, other.zw);
```

```
> error: no matching function for call to [...]
> template argument deduction/substitution failed: [...]
```

# swizzler<> PROBLEM

```
vec3 v = vec4(other.xy, other.zw);
```

```
> error: no matching function for call to [...]
> template argument deduction/substitution failed: [...]
```

- Solution? Introduce another abstraction layer!

# swizzler<> PROBLEM

```
vec3 v = vec4(other.xy, other.zw);
```

```
> error: no matching function for call to [...]
> template argument deduction/substitution failed: [...]
```

- Solution? Introduce another abstraction layer!
- `(construct_at_index(i, decay(std::forward<Args>(args))), ...);`

# swizzler<> PROBLEM

```
vec3 v = vec4(other.xy, other.zw);
```

```
> error: no matching function for call to [...]
> template argument deduction/substitution failed: [...]
```

- Solution? Introduce another abstraction layer!
- `(construct_at_index(i, decay(std::forward<Args>(args))), ...);`
- `decay` calls equivalent member function (or does nothing for scalar)

# swizzler<> PROBLEM

```
vec3 v = vec4(other.xy, other.zw);
```

```
> error: no matching function for call to [...]
> template argument deduction/substitution failed: [...]
```

- Solution? Introduce another abstraction layer!
- `(construct_at_index(i, decay(std::forward<Args>(args))), ...);`
- `decay` calls equivalent member function (or does nothing for scalar)
- both `vector` and `swizzler` have one so they can interchange easily

# OPERATORS AND FUNCTIONS

We will need to re-create a lot of generic utility functions

We will need to re-create a lot of generic utility functions

Example: the dot (inner) product of two vectors

We will need to re-create a lot of generic utility functions

Example: the dot (inner) product of two vectors

```
template<typename T, size_t... Ns>
T dot(const vector<T, Ns...> &, const vector<T, Ns...> &);

float n = dot(vec3(1, 0, 0), vec3(0, 0, 1));
```



We immediately hit a big problem!

```
vec3 v = vec3(1, 0, 0);
float n = dot(v.xzx, v.zyx);
```

```
> 'dot': no matching overloaded function found
> could not deduce template argument
```

- i Type deduction does not consider implicit conversions!

- **i** Type deduction does not consider implicit conversions!
  - Possible fixes:



Type deduction does not consider implicit conversions!

- Possible fixes:
- `float n = dot<float, 0, 1, 2>(v.xzx, v.zyx);`



Type deduction does not consider implicit conversions!

- Possible fixes:
- `float n = dot<float, 0, 1, 2>(v.xzx, v.zyx);`
- create by hand all scalar/size combinations :(



Type deduction does not consider implicit conversions!

- Possible fixes:
- `float n = dot<float, 0, 1, 2>(v.xzx, v.zyx);`
- create by hand all scalar/size combinations :(
- SFINAE tricks

## A BETTER FIX

- i we place the functions in a non-deduced context: inside `vector<>` itself!

# A BETTER FIX

- i we place the functions in a non-deduced context: inside `vector<>` itself!

```
template<typename T, size_t... Ns>
struct vector
{
    friend T dot(const vector& a, const vector& b)
    {
        /* inline friend found via ADL */
    }
    (etc)
```

# ARE WE DONE?

 No!

# ARE WE DONE?



No!

Lots of shader code uses scalar types only:

```
float opS(float d1, float d2)
{
    return max(-d2, d1);
}
```

# ARE WE DONE?



No!

Lots of shader code uses scalar types only:

```
float opS(float d1, float d2)
{
    return max(-d2, d1);
}
```

We only provide the vector variant

```
friend vector max(const vector& a, const vector& b)
{
    return vector((a.data[Ns] < b.data[Ns] ? a.data[Ns] : b.data[Ns]))
}
```

# MORE TROUBLE

# MORE TROUBLE

Ambiguity with literals:

```
smoothstep(0, 1, v.xyz);  
//           ^  
//           could be `int`, `float` or `double`
```

# MORE TROUBLE

Ambiguity with literals:

```
smoothstep(0, 1, v.xyz);  
//           ^  
//           could be `int`, `float` or `double`
```

```
friend vector smoothstep(scalar_type, scalar_type, const vector&);  
//           ^  
//           only templated on vector
```

# SOLUTION

Inspect the list of function args and deduct a vector type using `std::common_type` techniques

# SOLUTION

Inspect the list of function args and deduct a vector type using `std::common_type` techniques

First: need to make `vec1` convert to/from scalars

```
std::is_convertible<vec1, float>::value == true
```

# SOLUTION

Inspect the list of function args and deduct a vector type using `std::common_type` techniques

First: need to make `vec1` convert to/from scalars

```
std::is_convertible<vec1, float>::value == true
```

Then: provide custom type traits

```
promote_to_vec< float >::type == vec1
```

# SOLUTION

Inspect the list of function args and deduct a vector type using `std::common_type` techniques

First: need to make `vec1` convert to/from scalars

```
std::is_convertible<vec1, float>::value == true
```

Then: provide custom type traits

```
promote_to_vec< float >::type == vec1
```

```
promote_to_vec< vec3 >::type == vec3
```

```
promote_to_vec< decltype(vec3().xyz) >::type == vec3
```

# SOLUTION

Inspect the list of function args and deduct a vector type using `std::common_type` techniques

First: need to make `vec1` convert to/from scalars

```
std::is_convertible<vec1, float>::value == true
```

Then: provide custom type traits

```
promote_to_vec< float >::type == vec1
```

```
promote_to_vec< vec3 >::type == vec3
```

```
promote_to_vec< decltype(vec3().xyz) >::type == vec3
```

```
promote_to_vec< vec3, float >::type == vec3
```

```
promote_to_vec< vec3, float, double >::type == vec3
```

# SOLUTION (CONT.)

# SOLUTION (CONT.)

Instead of friend functions are static

```
template<template<class, size_t...> class vector, class T, size_t...>
struct builtin_func_lib
{
    static vector max(const vector& a, const vector& b)
(etc)
```

# SOLUTION (CONT.)

Instead of friend functions are static

```
template<template<class, size_t...> class vector, class T, size_t...>
struct builtin_func_lib
{
    static vector max(const vector& a, const vector& b)
(etc)
```

Create an all-forwarding monster function

```
template<class... Args>
inline auto func(Args&&... args) ->
    decltype(decay(
        promote_to_vec<Args...>::type::
        func(std::forward<Args>(args)...)))
{
    return
        promote_to_vec<Args...>::type::
        func(std::forward<Args>(args)...);
}
```

# SOLUTION - COMPROMISE

```
...  
MAKE_LIB_FUNC(abs)  
MAKE_LIB_FUNC(sign)  
MAKE_LIB_FUNC(floor)  
MAKE_LIB_FUNC(trunc)  
MAKE_LIB_FUNC(ceil)  
MAKE_LIB_FUNC(fract)  
MAKE_LIB_FUNC(mod)  
MAKE_LIB_FUNC(min)  
MAKE_LIB_FUNC(max)  
MAKE_LIB_FUNC(clamp)  
MAKE_LIB_FUNC(mix)  
MAKE_LIB_FUNC(step)  
MAKE_LIB_FUNC(smoothstep)  
...
```

# SOLUTION - DEMO

```
return max(-d2, d1);
//           ||
//           ||
//           \
promote_to_vec<float, float>
//           ||
//           ||
//           \
return vector<float, 0>::max(-d2, d1);
```



## THE `matrix<>` DATATYPE

Now that we have `vector<>` a matrix is more straightforward

```
template<
    typename,
    template<typename, size_t...> class vector_type,
    typename...
>
struct matrix;
```

## FOLDING HELPER

- i two dimensional → introduce helper type for the indices

```
template<size_t...>
struct indices_pack;
```

# THE matrix<> DATATYPE (CONT.)

# THE matrix<> DATATYPE (CONT.)

```
template<
    typename scalar_type,
    template<typename, size_t...> class vector_type,
    size_t... Columns,
    size_t... Rows
>
struct matrix<scalar_type, vector_type,
    indices_pack<Columns...>, indices_pack<Rows...>>
```

# THE `matrix<>` DATATYPE (CONT.)

```
template<
    typename scalar_type,
    template<typename, size_t...> class vector_type,
    size_t... Columns,
    size_t... Rows
>
struct matrix<scalar_type, vector_type,
    indices_pack<Columns...>, indices_pack<Rows...>>

{
    static constexpr auto N = sizeof...(Columns);
    static constexpr auto M = sizeof...(Rows);

    using column_type = vector_type<scalar_type, Columns...>;
    using row_type = vector_type<scalar_type, Rows...>;

    column_type data[M];
    (etc)
```

# THE `matrix<>` DECLARATION

```
using vec2 = vector<float, 0, 1>;
using vec3 = vector<float, 0, 1, 2>;

using mat2 = matrix<float, vector,
    indices_pack<0, 1>, indices_pack<0, 1>>;
    
using mat3 = matrix<float, vector,
    indices_pack<0, 1, 2>, indices_pack<0, 1, 2>>; 

using mat2x3 = matrix<float, vector,
    indices_pack<0, 1>, indices_pack<0, 1, 2>>;
```

# THE `matrix<>` CONSTRUCTORS

# THE `matrix<>` CONSTRUCTORS

```
matrix() = default; // zeroes all data
```

# THE `matrix<>` CONSTRUCTORS

```
matrix() = default; // zeroes all data

explicit matrix(scalar_type s) // fill in diagonally
{
    ((data[Rows][Rows] = s), ...);
}
```

# THE `matrix<>` CONSTRUCTORS

```
matrix() = default; // zeroes all data

explicit matrix(scalar_type s) // fill in diagonally
{
    ((data[Rows][Rows] = s), ...);

template<typename... Args>
explicit matrix(Args&&... args)
{
    size_t i = 0;
    (construct_at_index(i,
        decay(std::forward<Args>(args))), ...);
}
```

## THE `matrix<>` OPS AND FUNCS

- can recycle the same binary operators as vector if written generic
- except *multiplication*
  - which needs to be handled differently
  - for all variations of `matrix`, `row_type`, `column_type`

## PRIOR ART

 not invented here :)

- clang vector extensions

- ```
typedef float vec3 __attribute__((ext_vector_type(3)));
```
- PRO: supports full swizzling
- CON: very limited in initializations

- 3rd party libraries
  - GLM
    - .xyz() style only
    - horrible preprocessor heavy implementation
  - CXXSwizzle
    - full spec
    - slow debug

# RESULTS

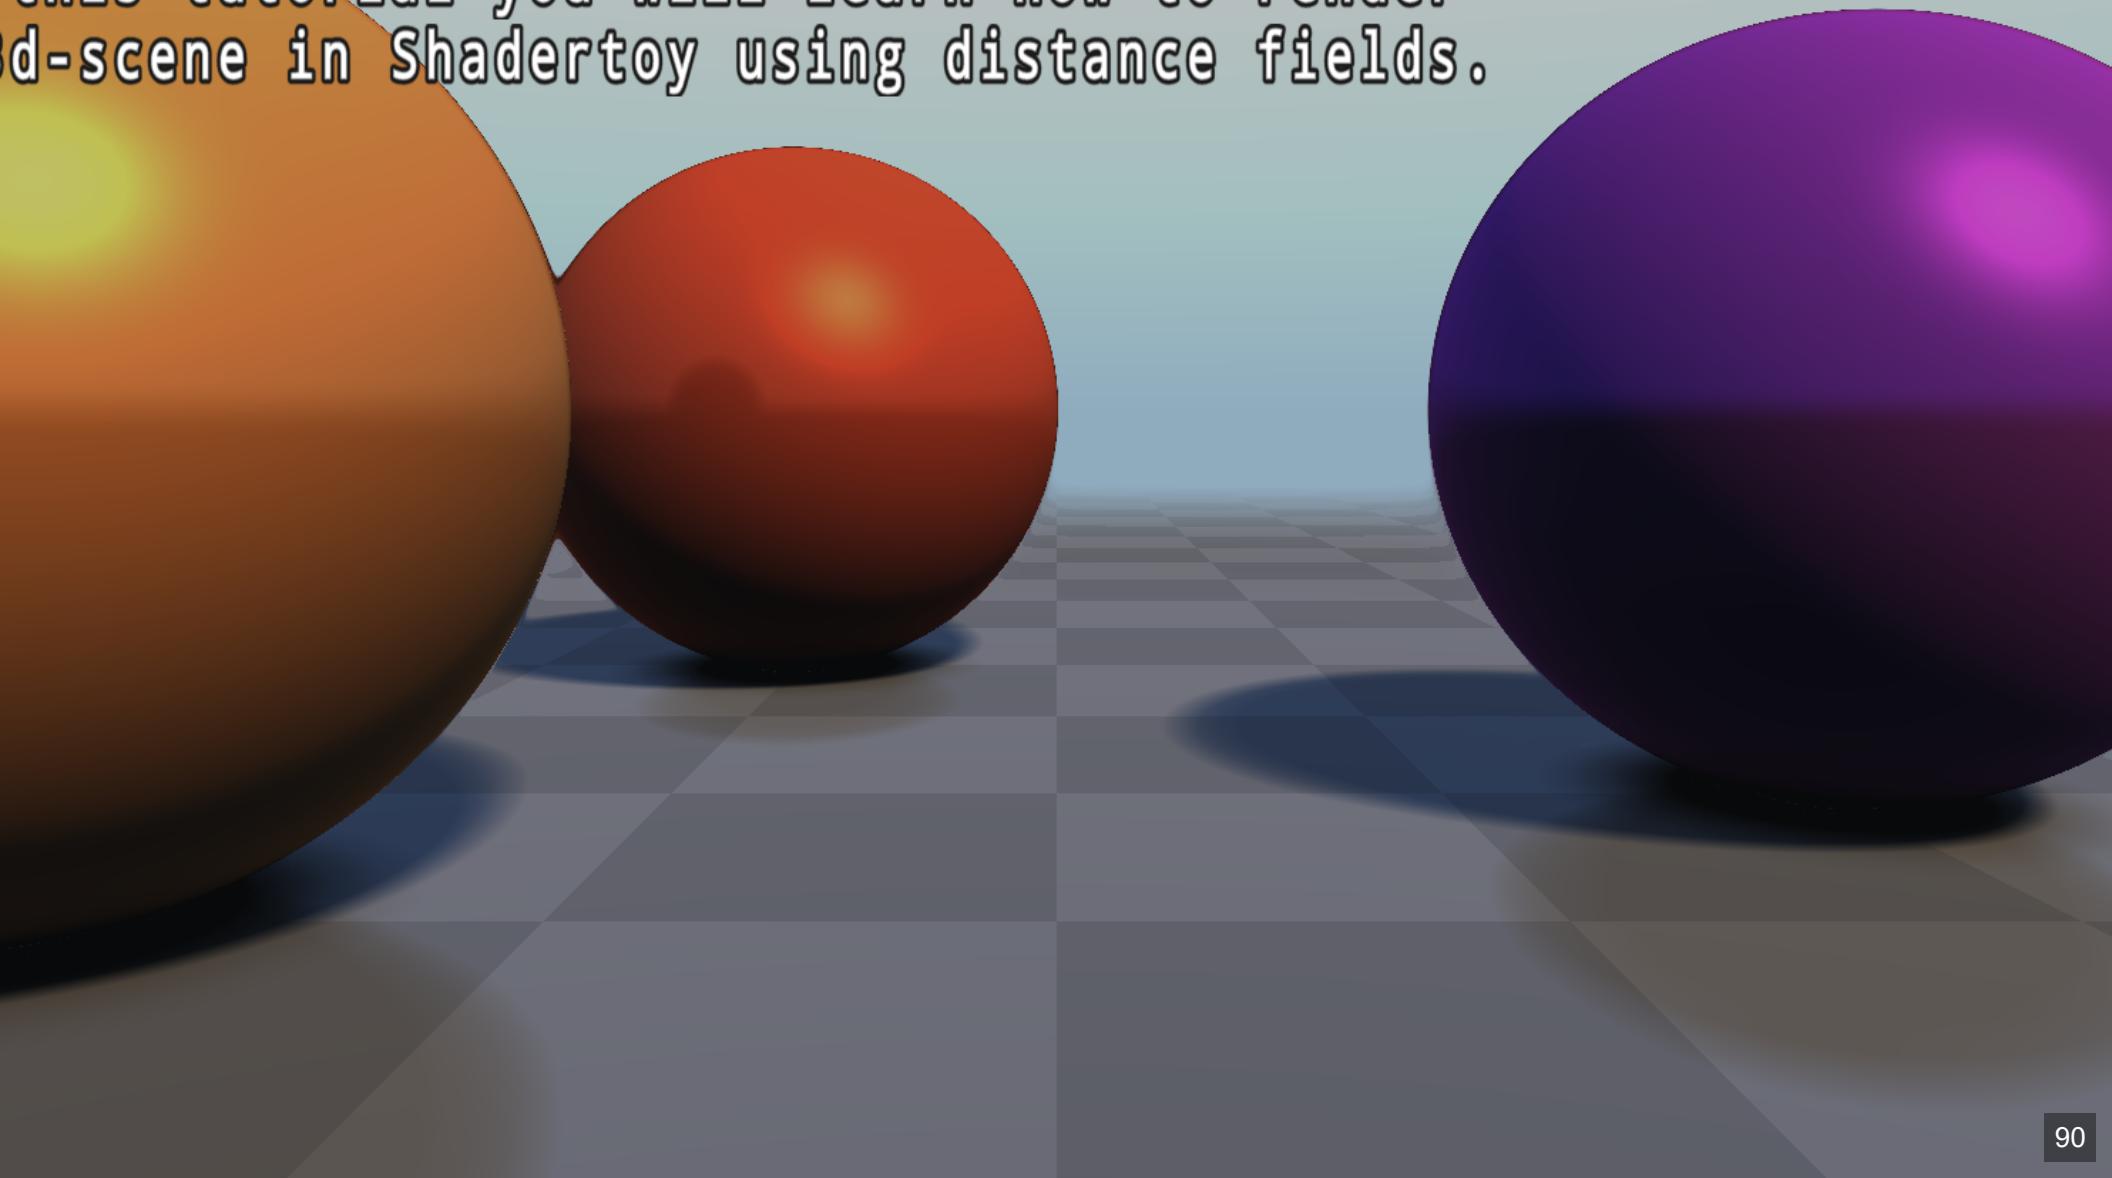
...but first: Crash Course into Procedural Graphics!

 Courtesy of @ReinderNijhoff

<https://www.shadertoy.com/view/4dSfRc>

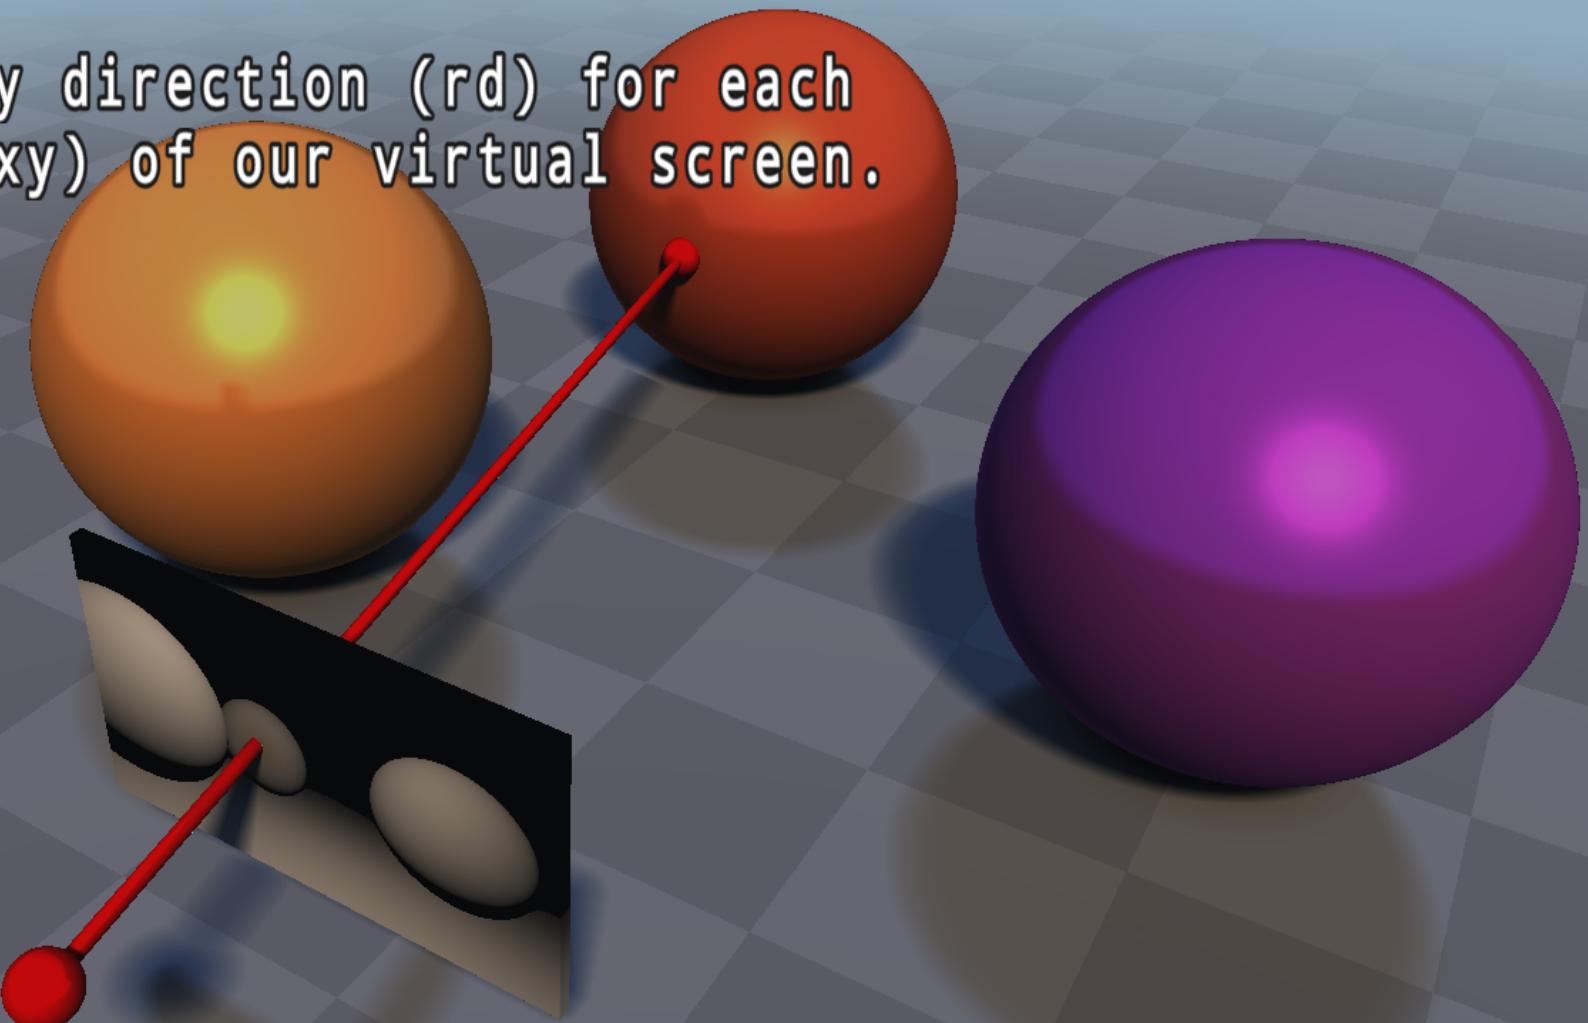
# Raymarching distance fields

In this tutorial you will learn how to render a 3d-scene in Shadertoy using distance fields.



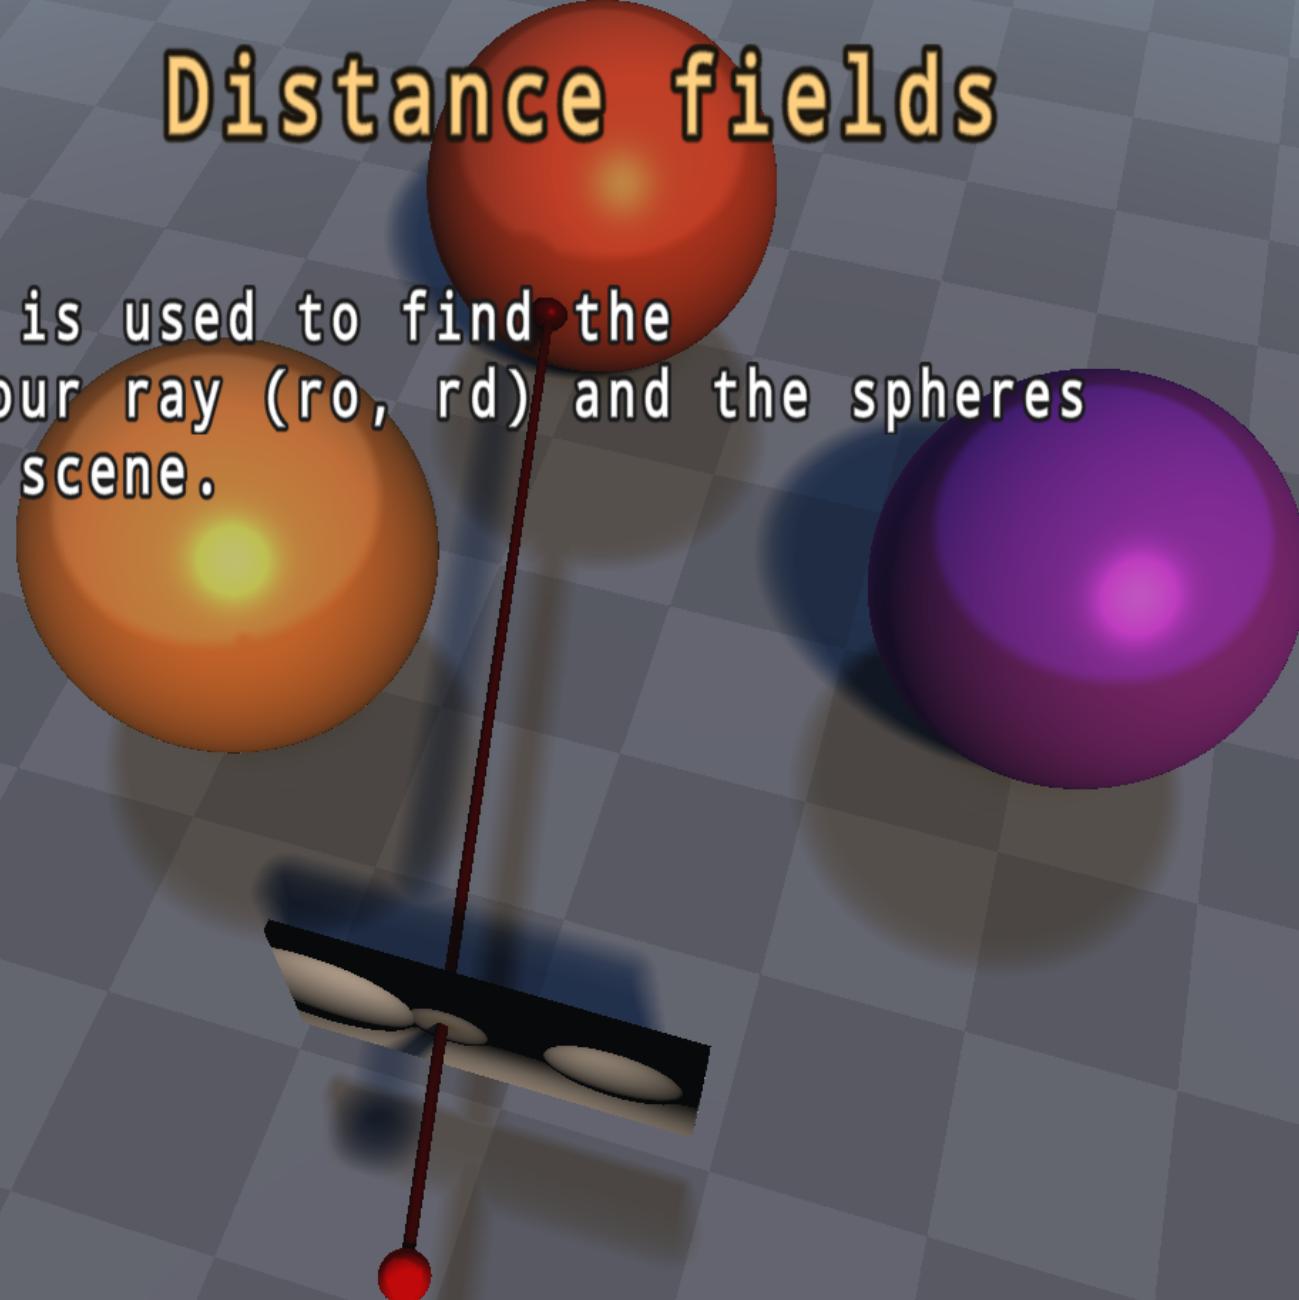
# Create a ray

compute the ray direction ( $rd$ ) for each pixel ( $fragCoord.xy$ ) of our virtual screen.



# Distance fields

A distance field is used to find the intersection of our ray ( $ro$ ,  $rd$ ) and the spheres in the plane of the scene.



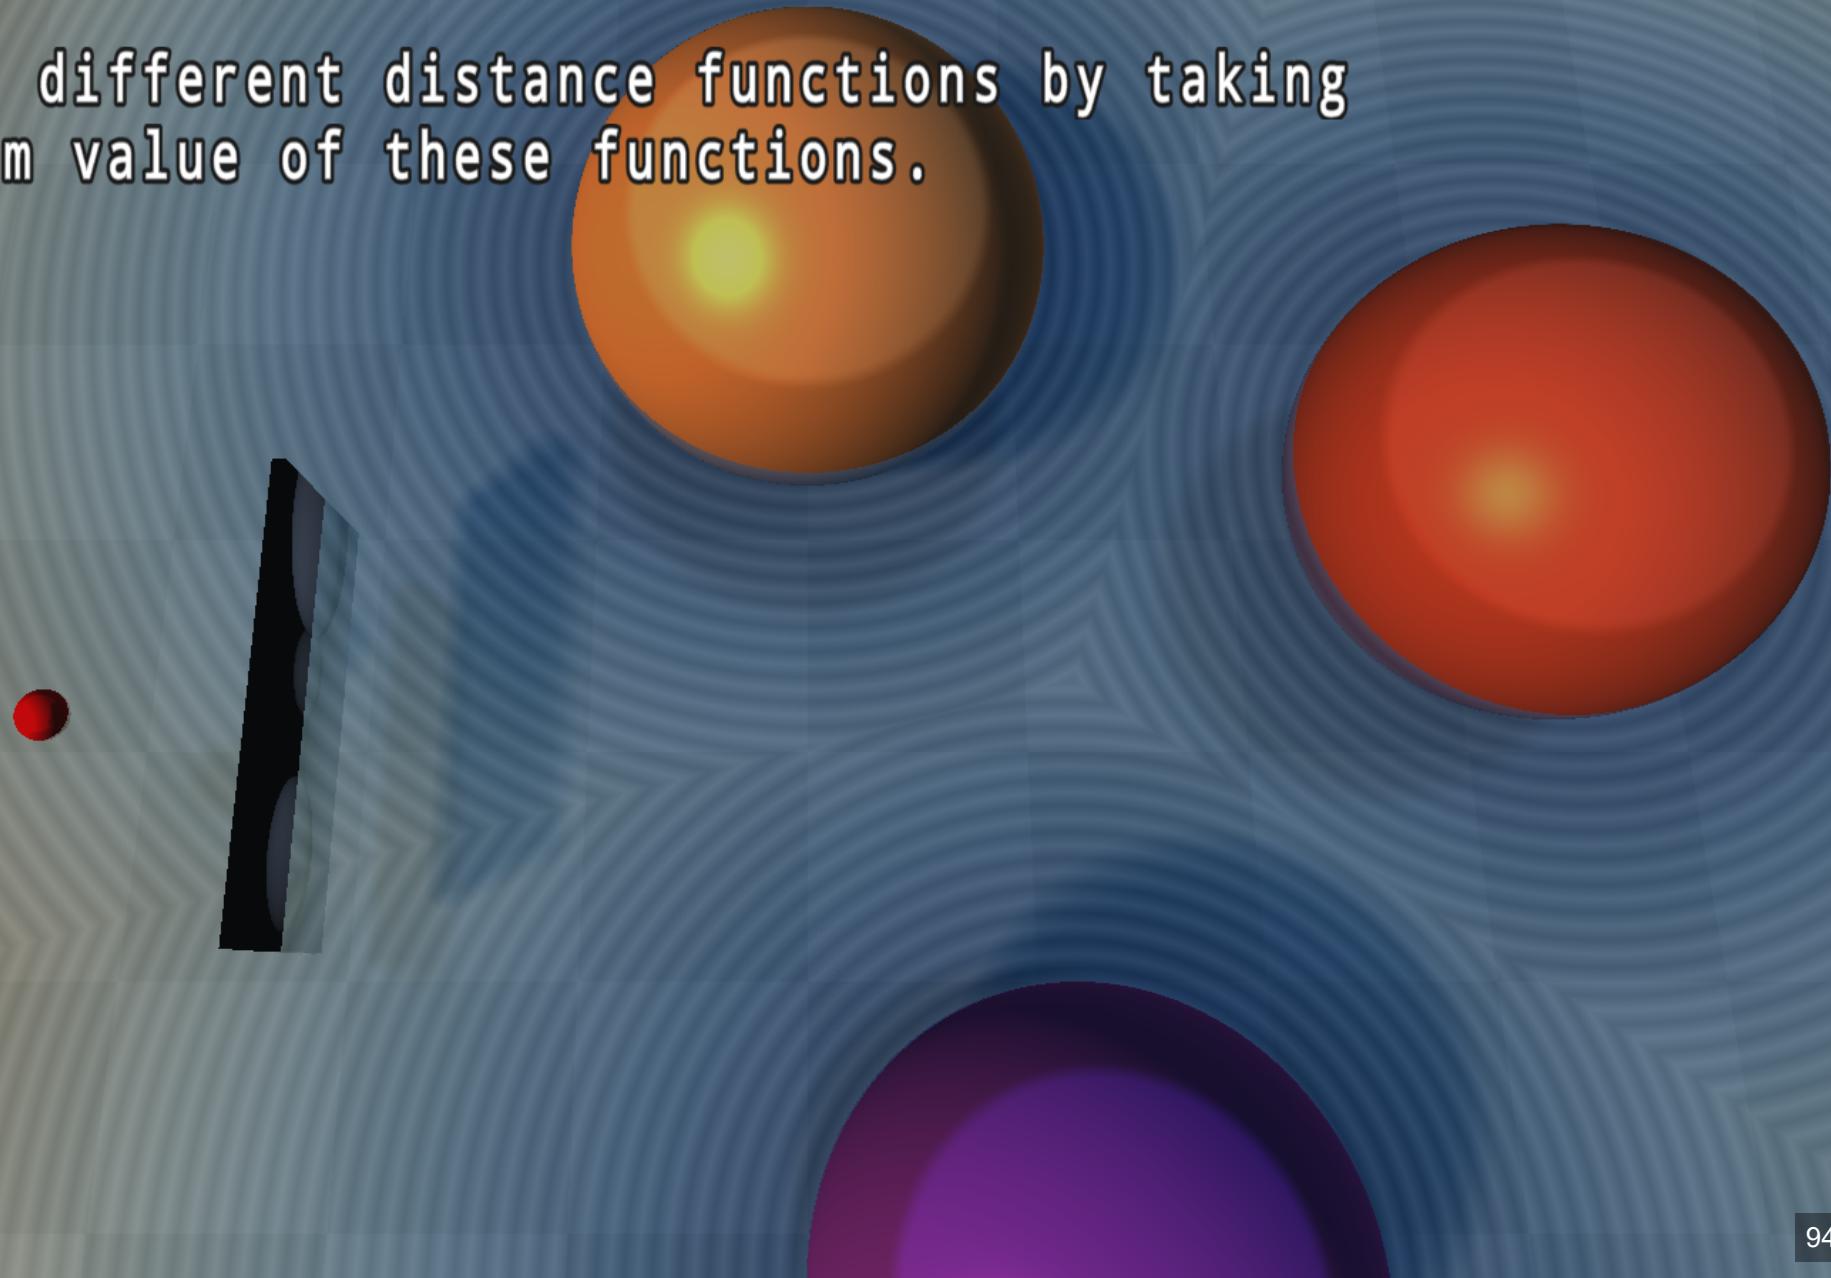
# Distance fields

The distance function for a sphere is the distance to center of the sphere minus the radius of the sphere.

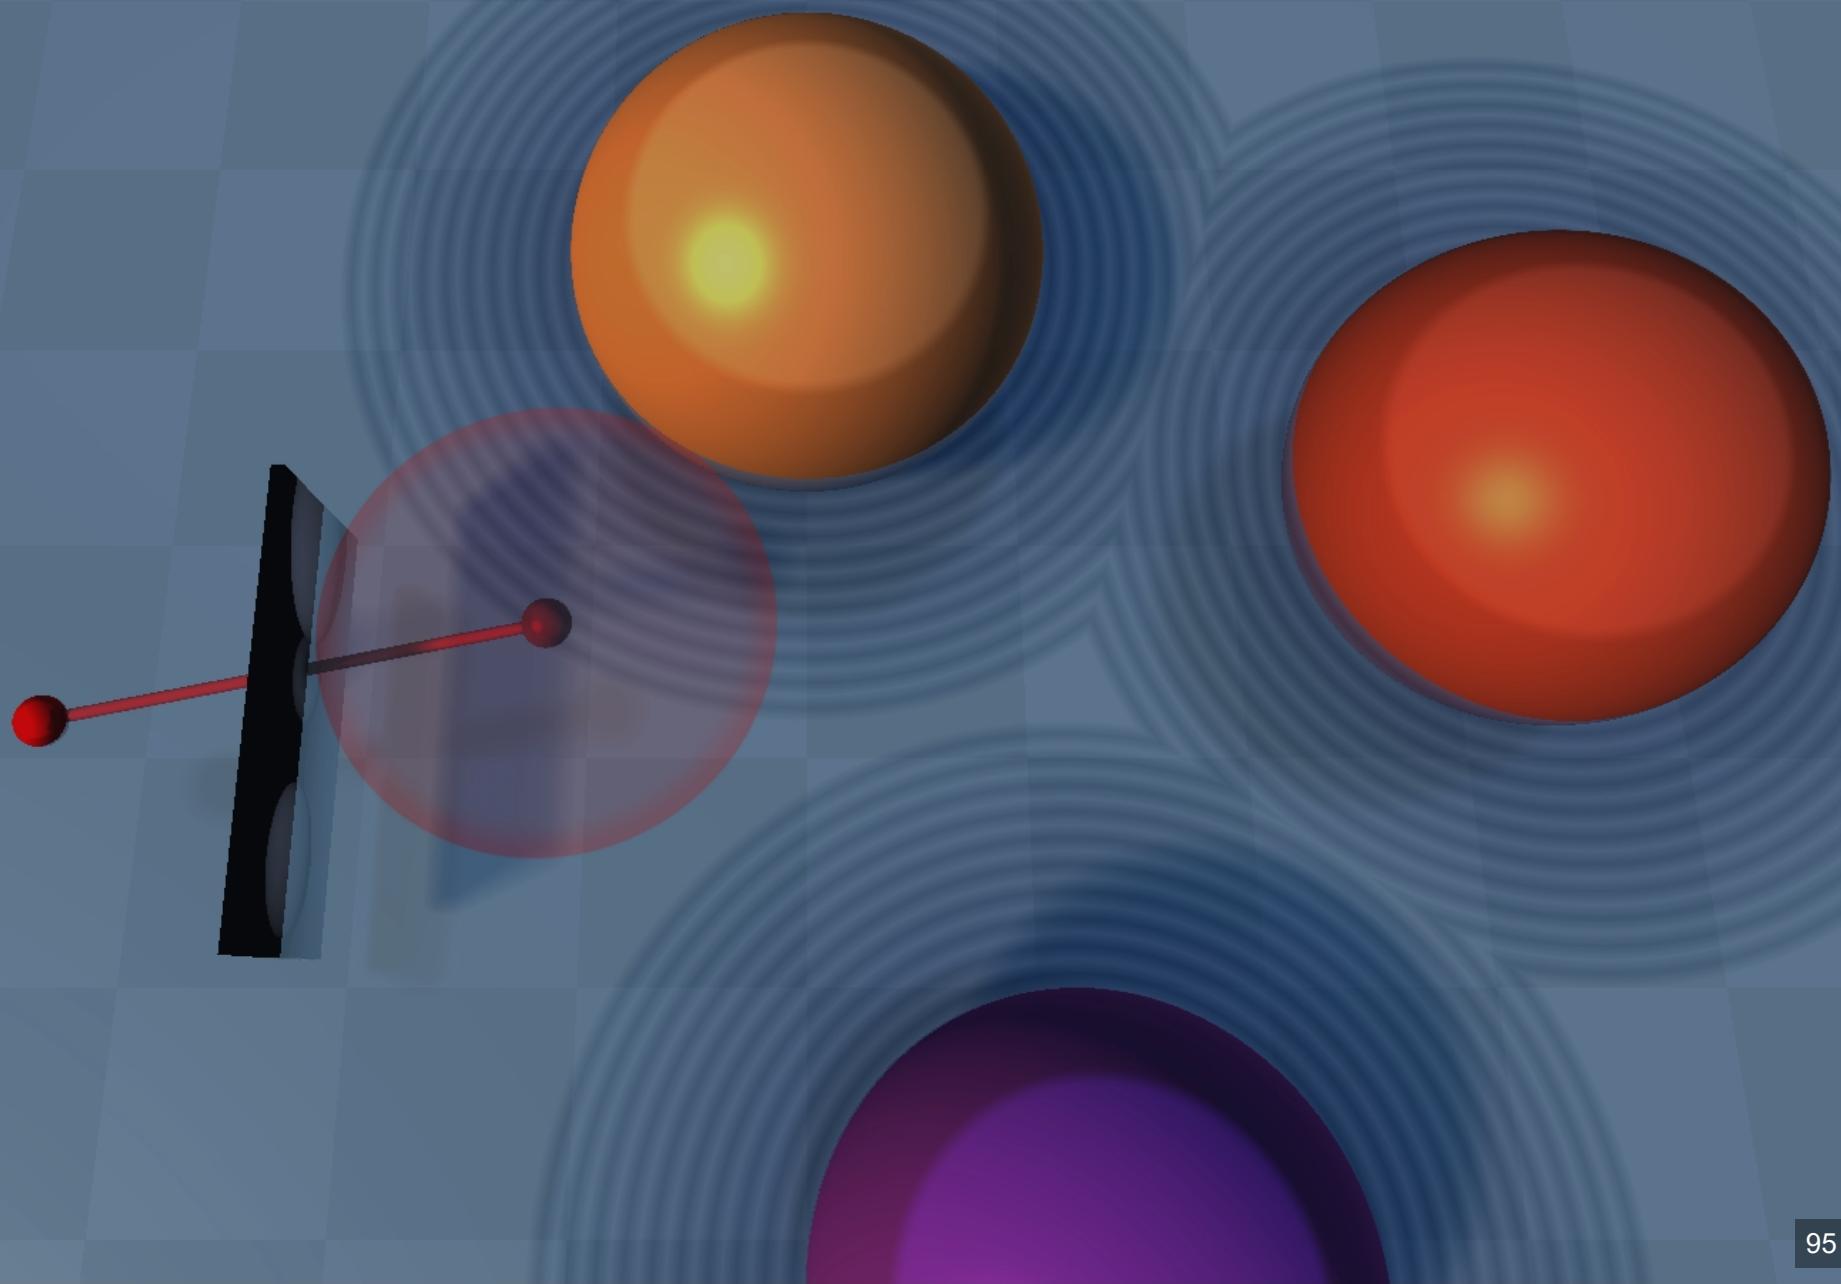


# Distance fields

combine different distance functions by taking minimum value of these functions.



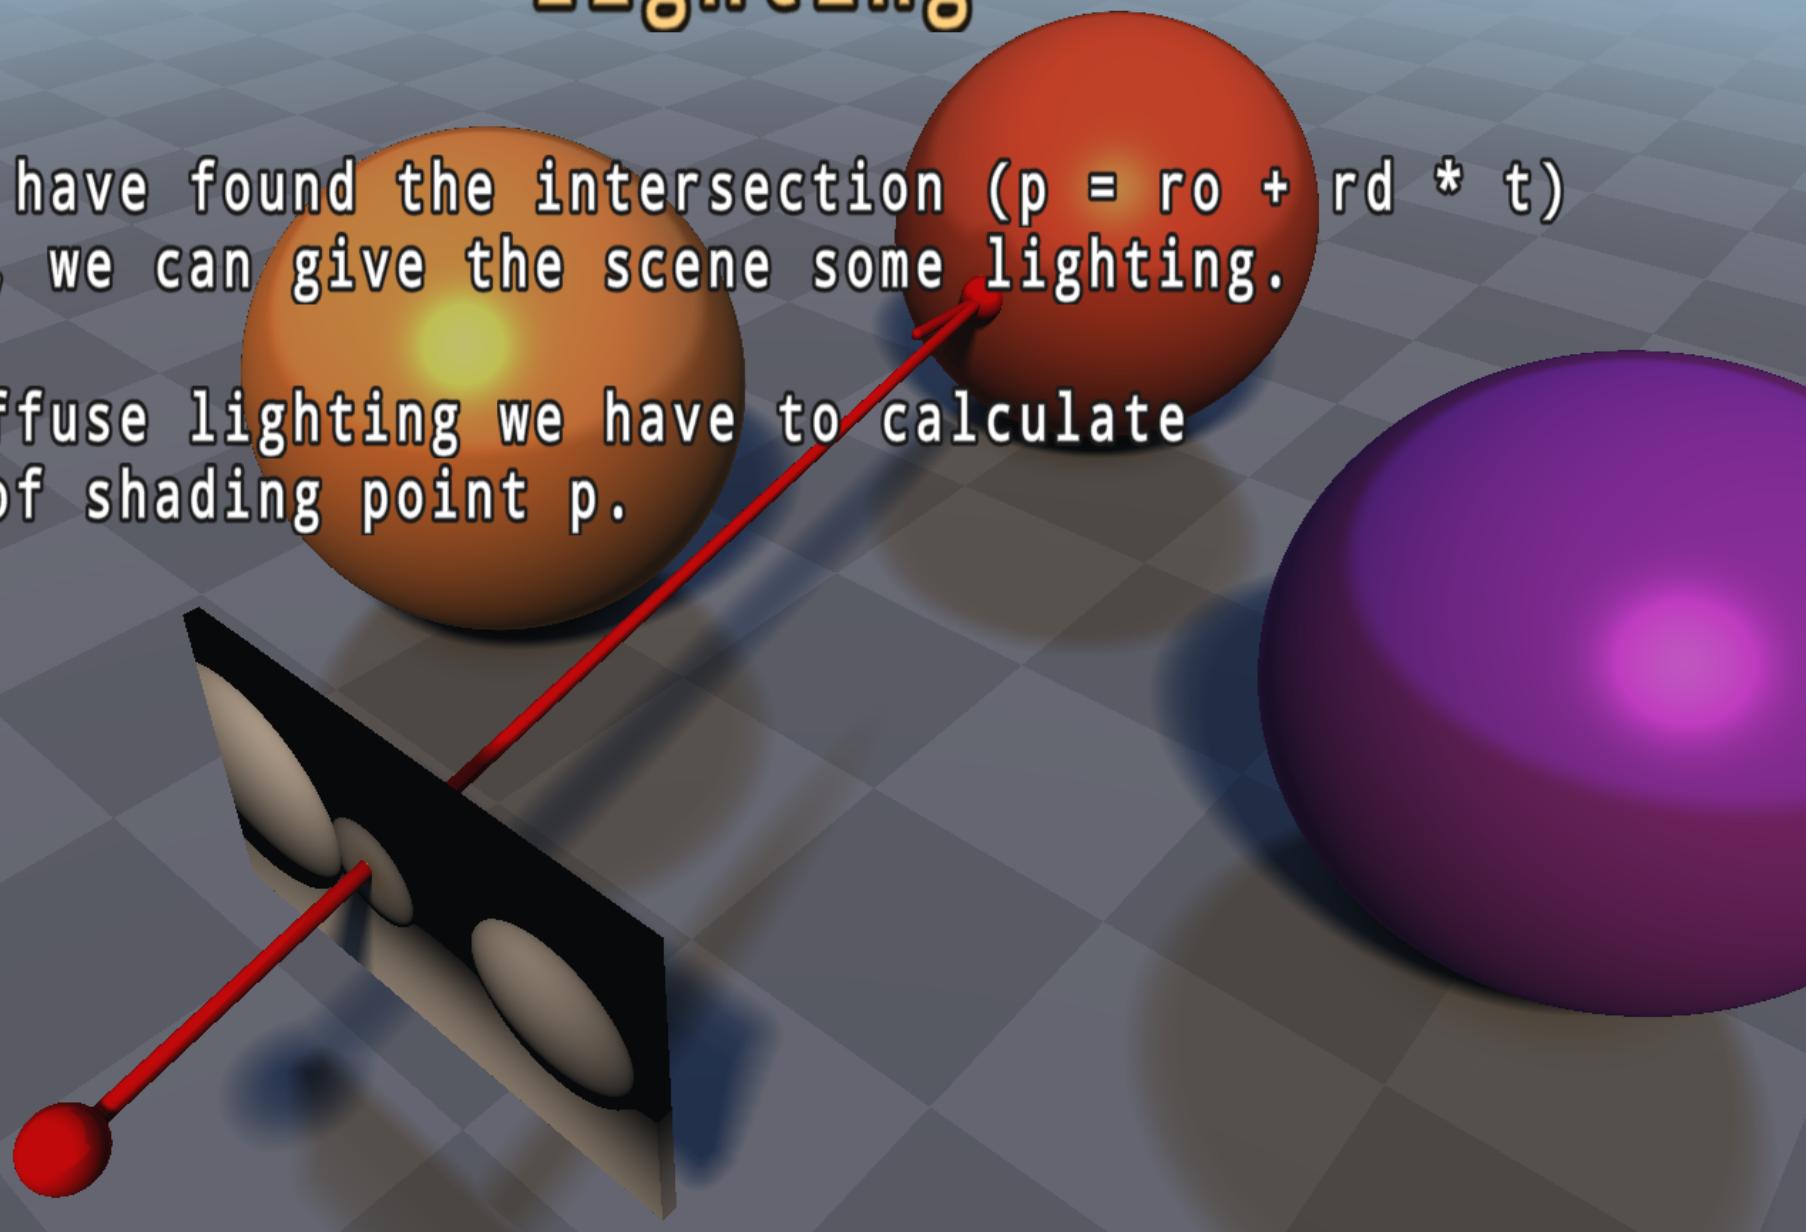
# Distance fields



# Lighting

Now that we have found the intersection ( $p = ro + rd * t$ ) of our ray, we can give the scene some lighting.

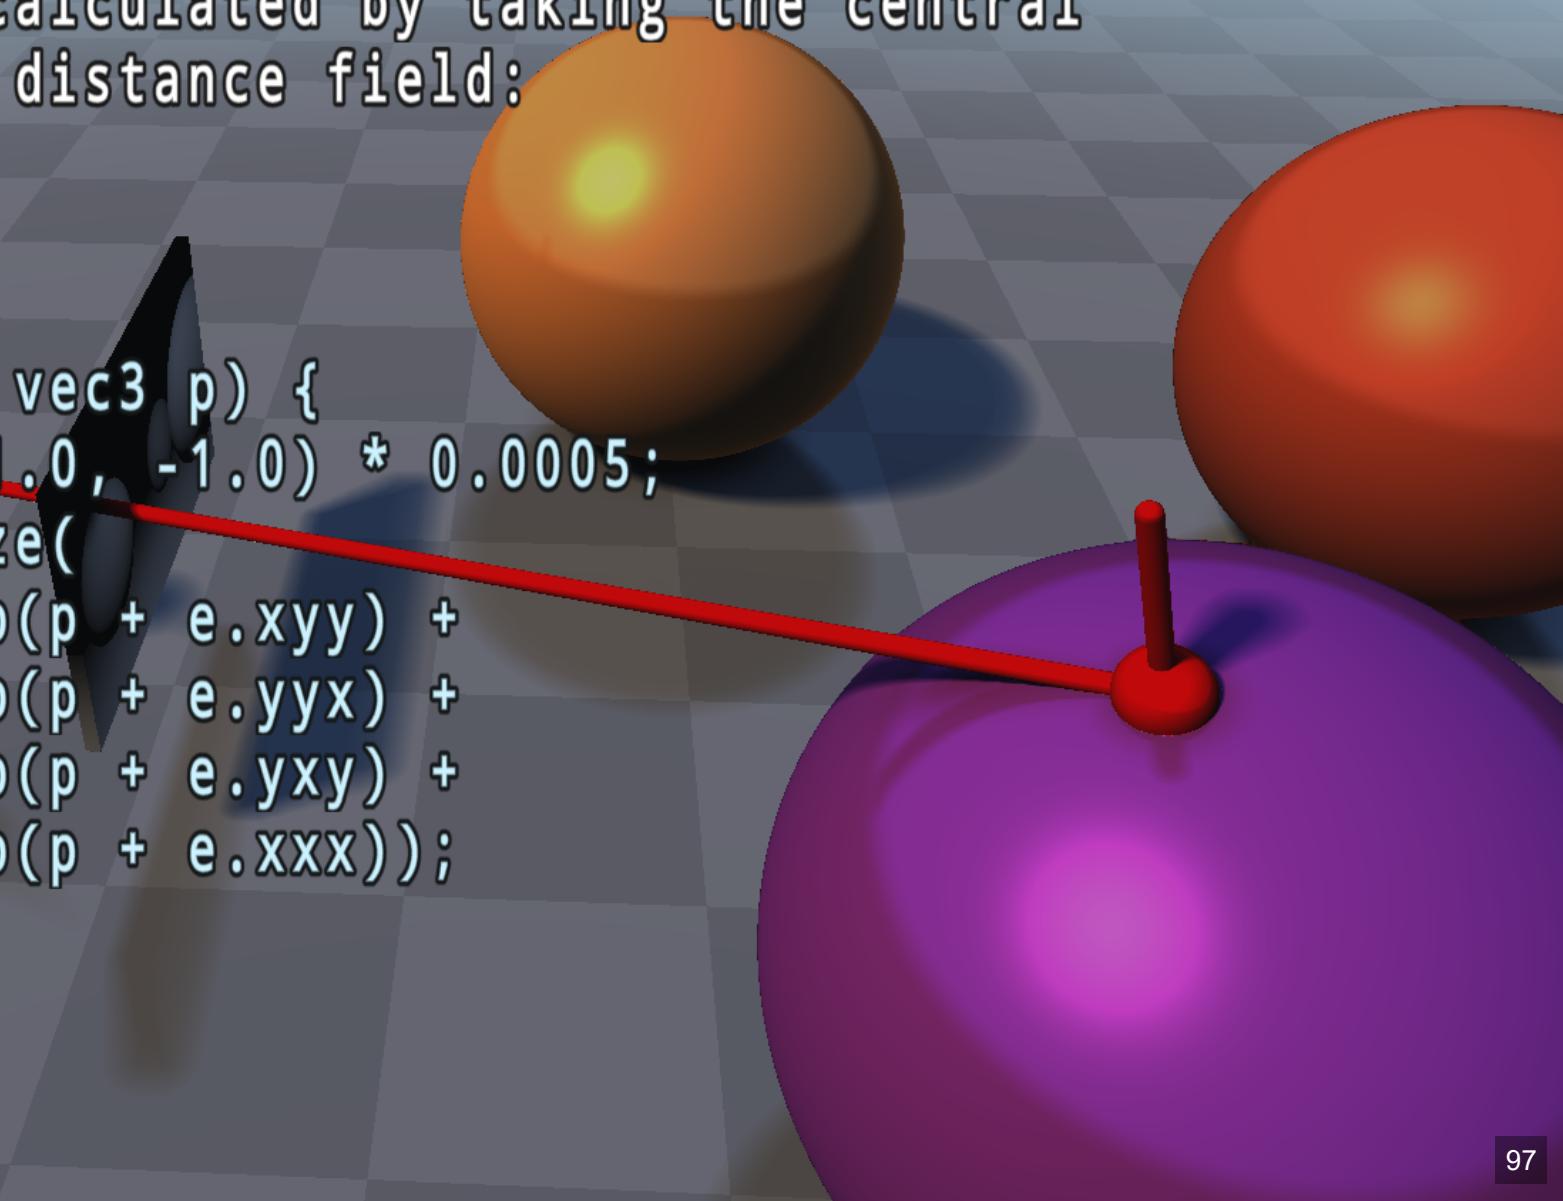
To apply diffuse lighting we have to calculate the normal of shading point p.



# Lighting

The normal can be calculated by taking the central differences on the distance field:

```
3 calcNormal(in vec3 p) {  
vec2 e = vec2(1.0, -1.0) * 0.0005;  
return normalize(  
    e.xyy * map(p + e.xyy) +  
    e.yyx * map(p + e.yyx) +  
    e.yxy * map(p + e.yxy) +  
    e.hxx * map(p + e.hxx));
```



# SHOWCASE

## GPU / DESKTOP PC

- <https://www.shadertoy.com/user/valentingalea>
- Nvidia GeForce 1060
- 1080p

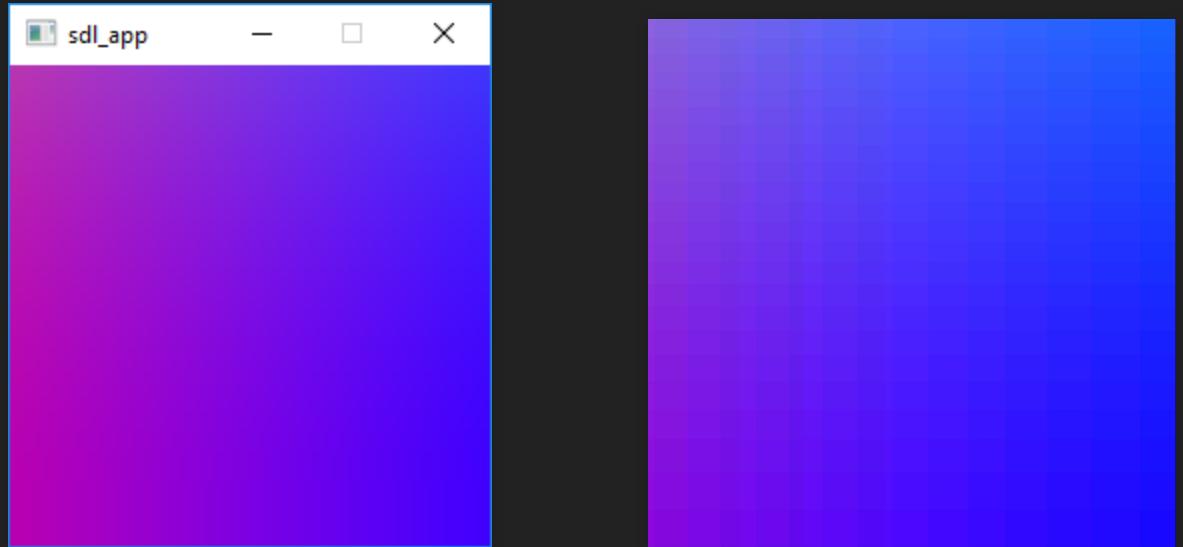
## CPU / DESKTOP PC

- minimal draw app with <https://www.libsdl.org/>
- AMD FX 8350 8-core 4.00 GHz
- Microsoft Visual C++ 2017
  - /O2 /Ob2 /fp:fast /fp:except-

## CPU / MOBILE PHONE

- C4Droid app  
(<https://play.google.com/store/apps/details?id=com.n0n3m4.droidc>)
- Samsung Galaxy S7
- GCC 8.0
  - -Ofast -march=native -funroll-loops

# HELLO WORLD (CPU)



---

240x240 px

85.62 FPS

166.77 FPS

---

120x120 px

100.27 FPS

468.49 FPS



# PLANET (CPU)



---

240x240 px

1.92 FPS

0.83 FPS

---

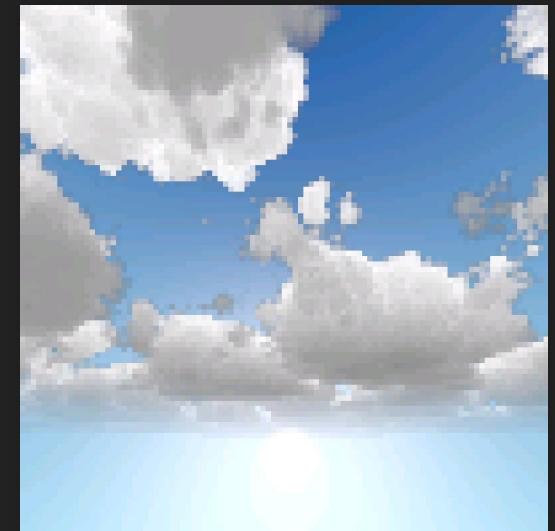
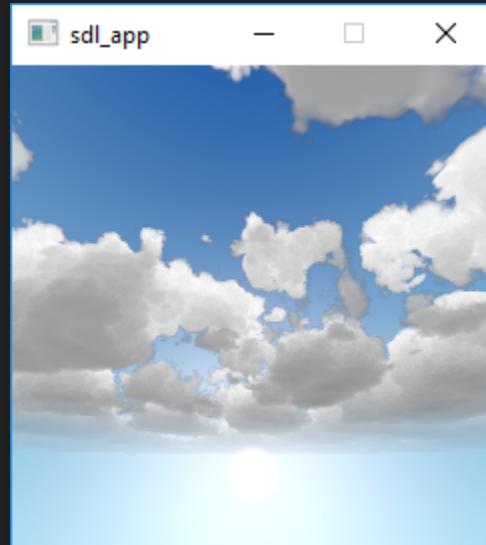
120x120 px

7.30 FPS

3.34 FPS



# CLOUDS (CPU)



---

240x240 px

2.54 FPS

2.44 FPS

---

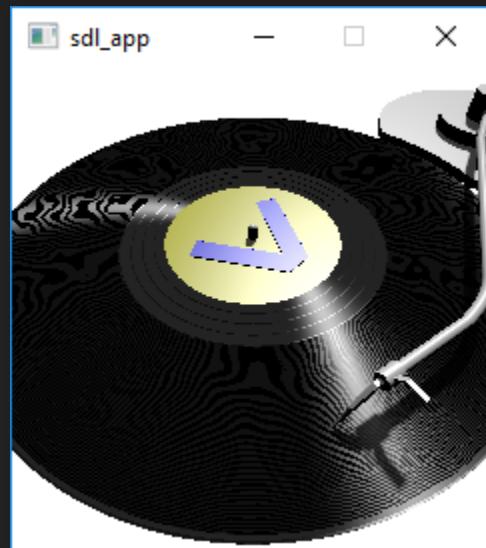
120x120 px

9.63 FPS

9.64 FPS



# VINYL TURNTABLE (CPU)



---

240x240 px

8.44 FPS

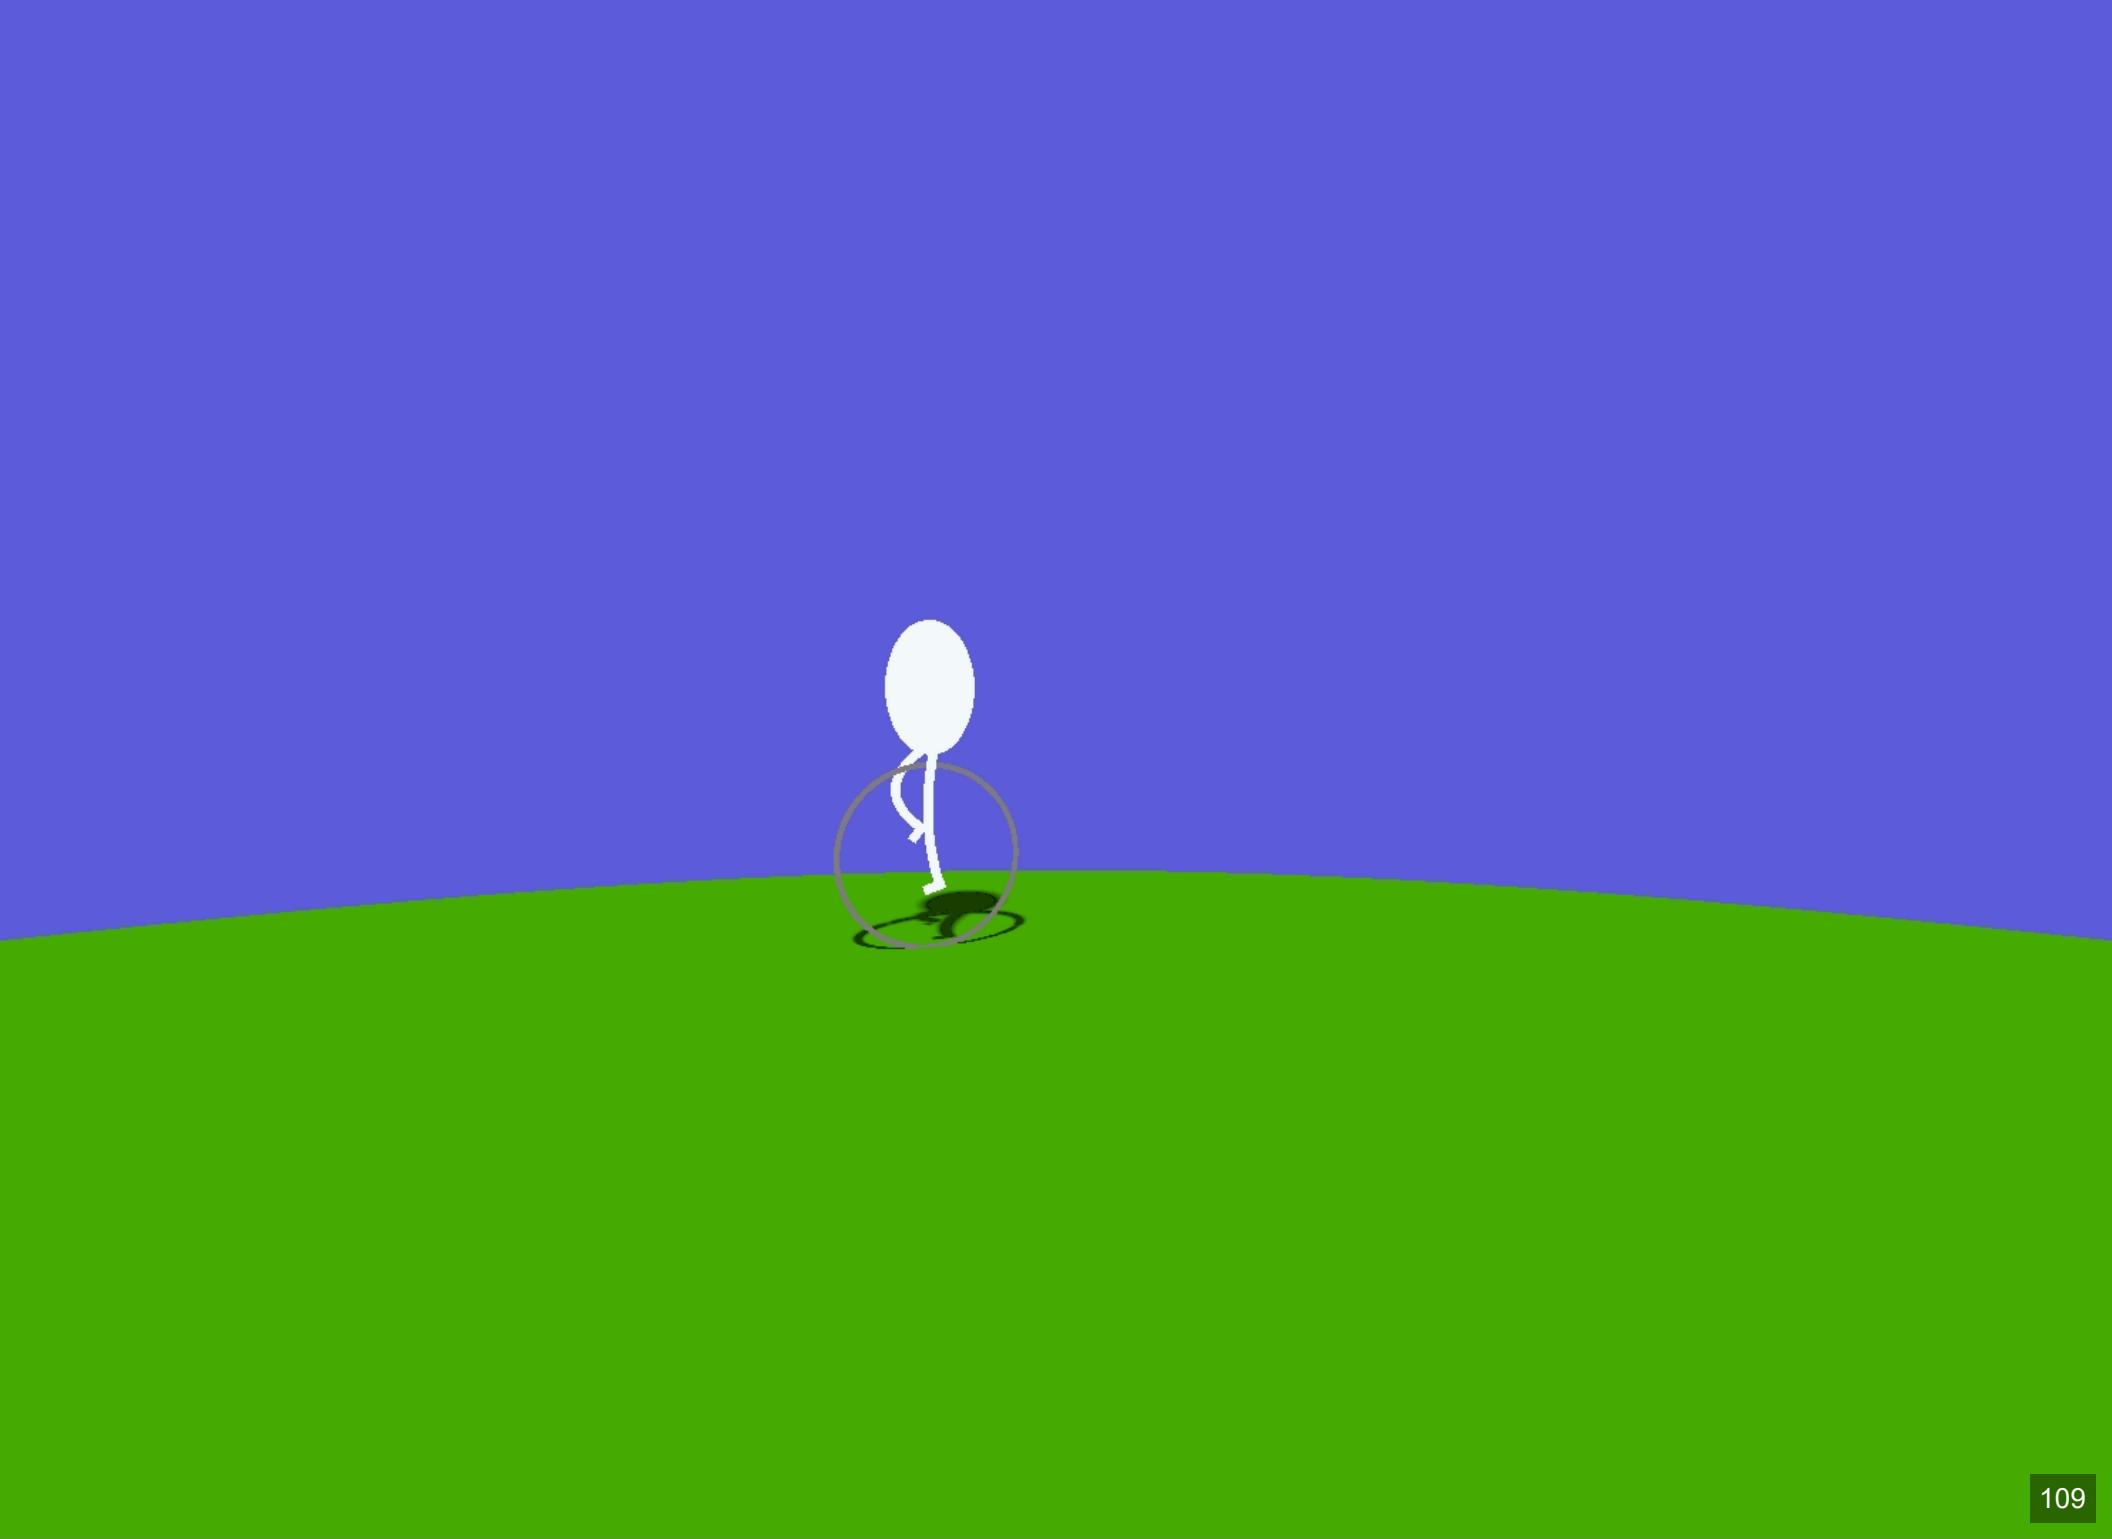
2.94 FPS

---

120x120 px

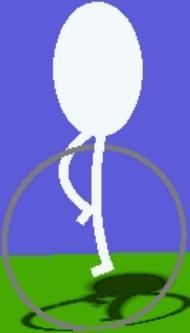
28.11 FPS

12.82 FPS





@valentin\_galea





@valentin\_galea



<https://github.com/valentingalea/>





@valentin\_galea



<https://github.com/valentingalea/>



<https://www.splashdamage.com/>

## ATTRIBUTION

Piotr Gwiazdowski @gwiazdorrr for original  
inspiration and help

Shading and Renderman: Jaume Sanchez |  
@thespite

Motivation Shaders: Inigo Quilez

<https://www.shadertoy.com/view/ld3Gz2>

<https://www.shadertoy.com/view/ldScDh>

<https://www.shadertoy.com/view/4ttSWf>

GPU pipeline: <https://open.gl/> (CC BY-SA 4.0)

All other images under "Fair Use"/"Fair Dealing"